

AutoTVM & AutoScheduler

Presenter : Jaehun Ryu

jaehunryu@postech.ac.kr

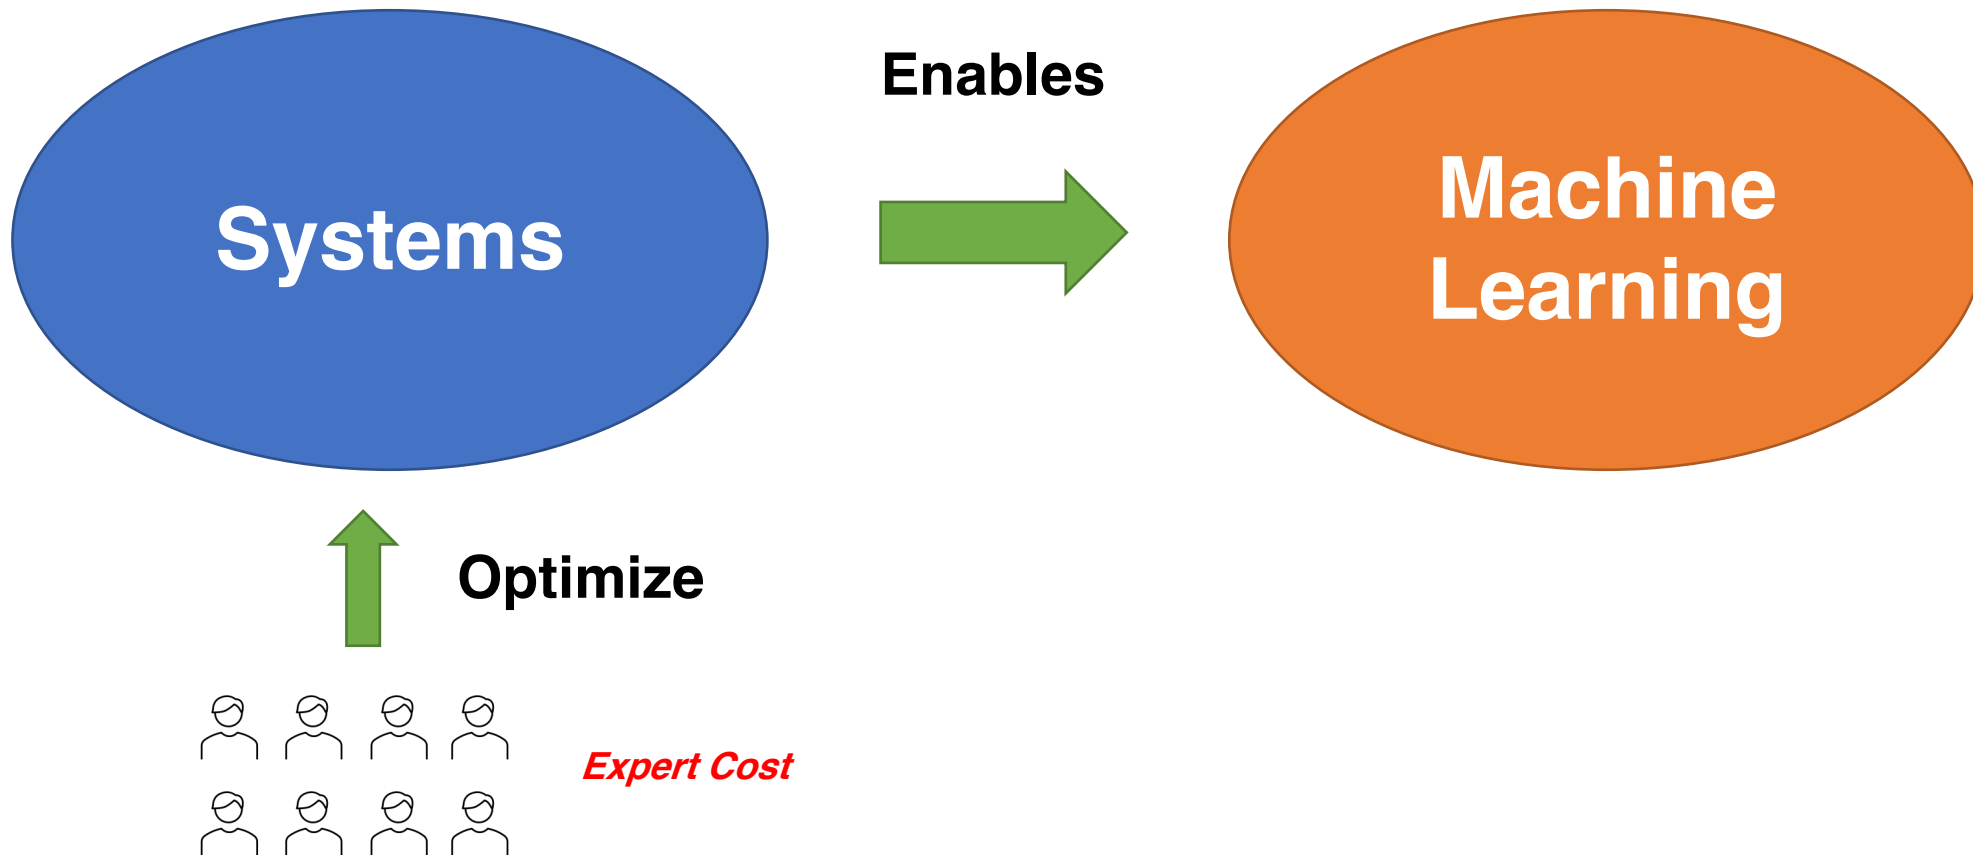
Introduction

- AutoTVM : Template-based Auto Tuning
 - Learning to optimize tensor programs(NIPS18,Chen et al)
- AutoScheduler : Template-free Auto Scheduling
 - Ansor: Generating High-Performance Tensor Programs for Deep Learning(OSDI 20,Zheng et al)

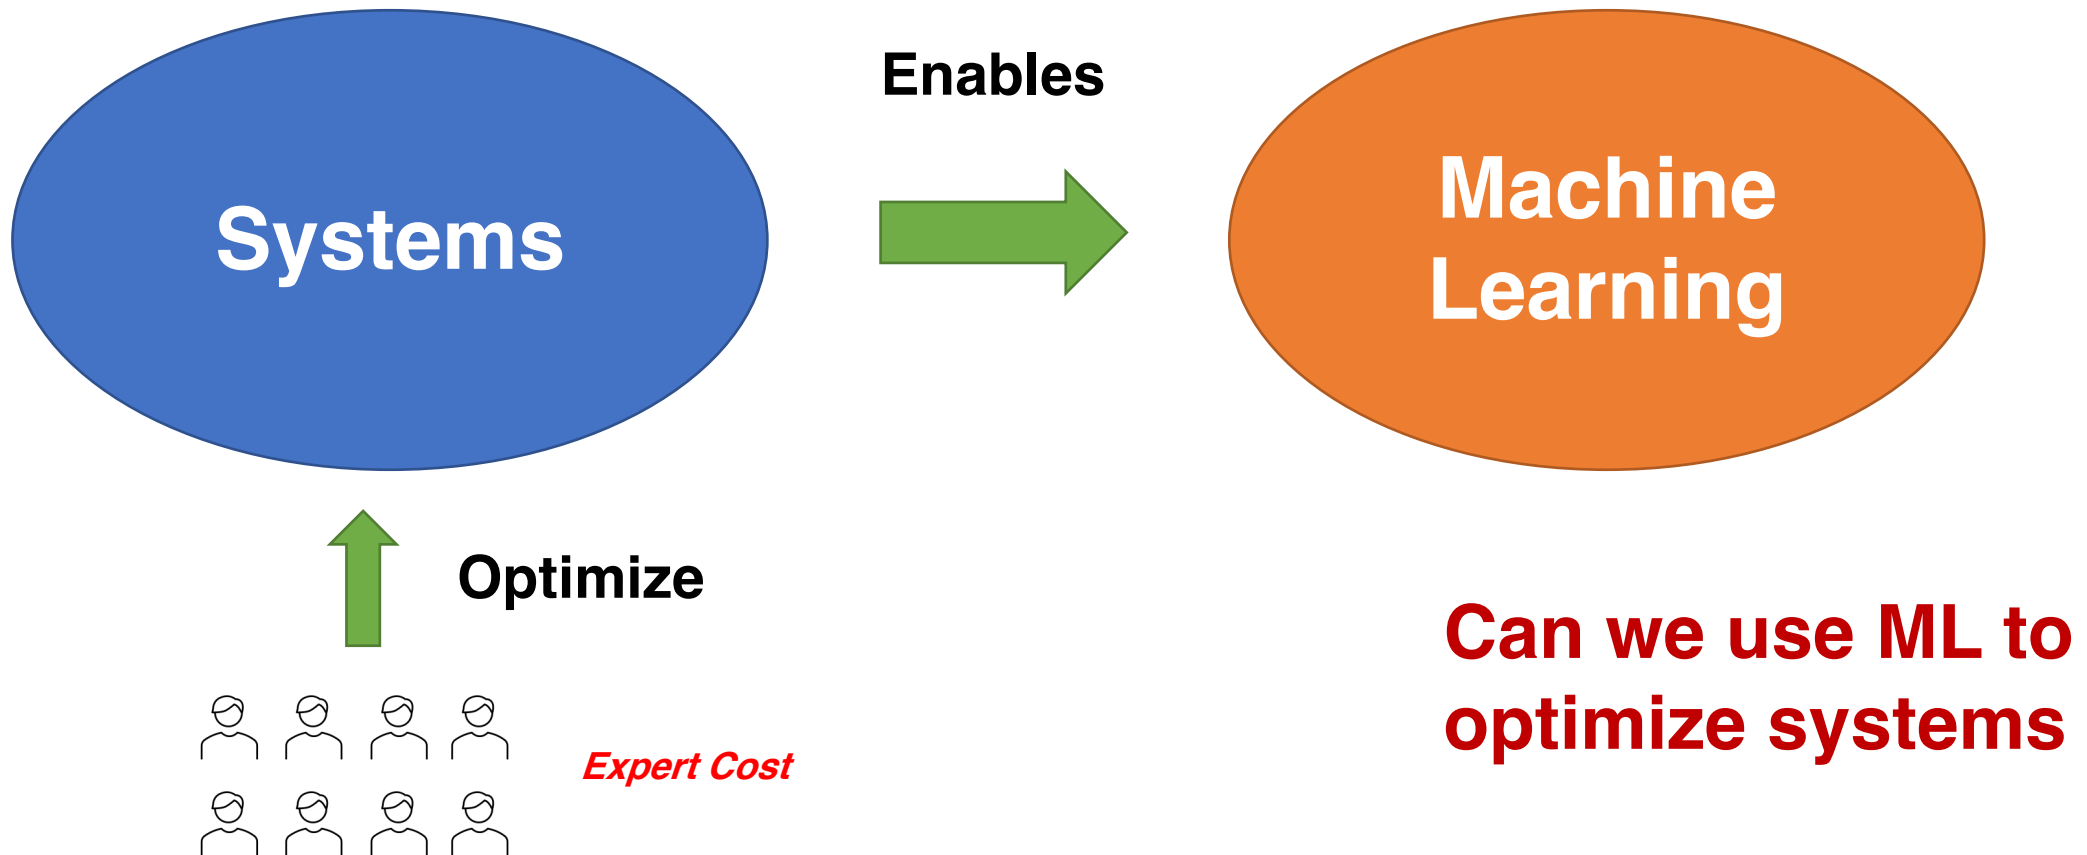
Both are TVM built-in autotuning methods.

Autotuning

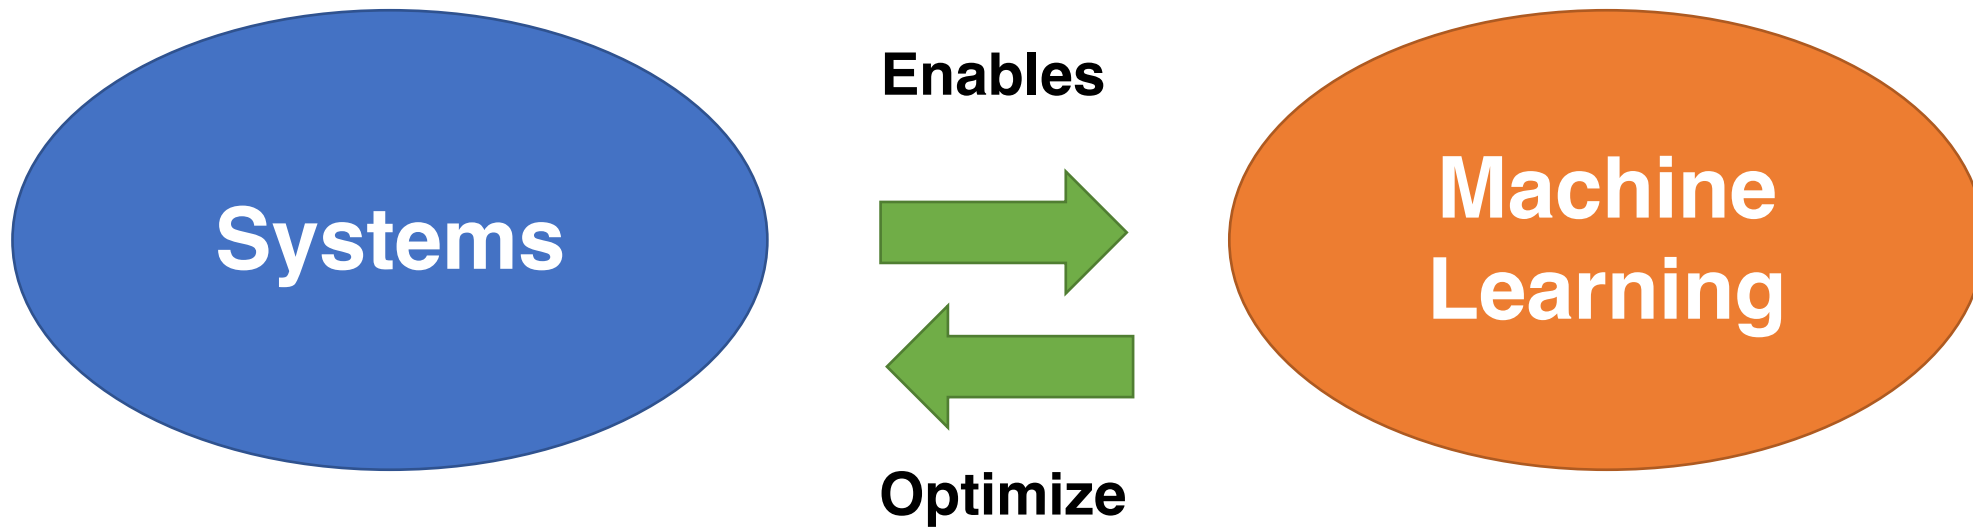
Current Learning Systems



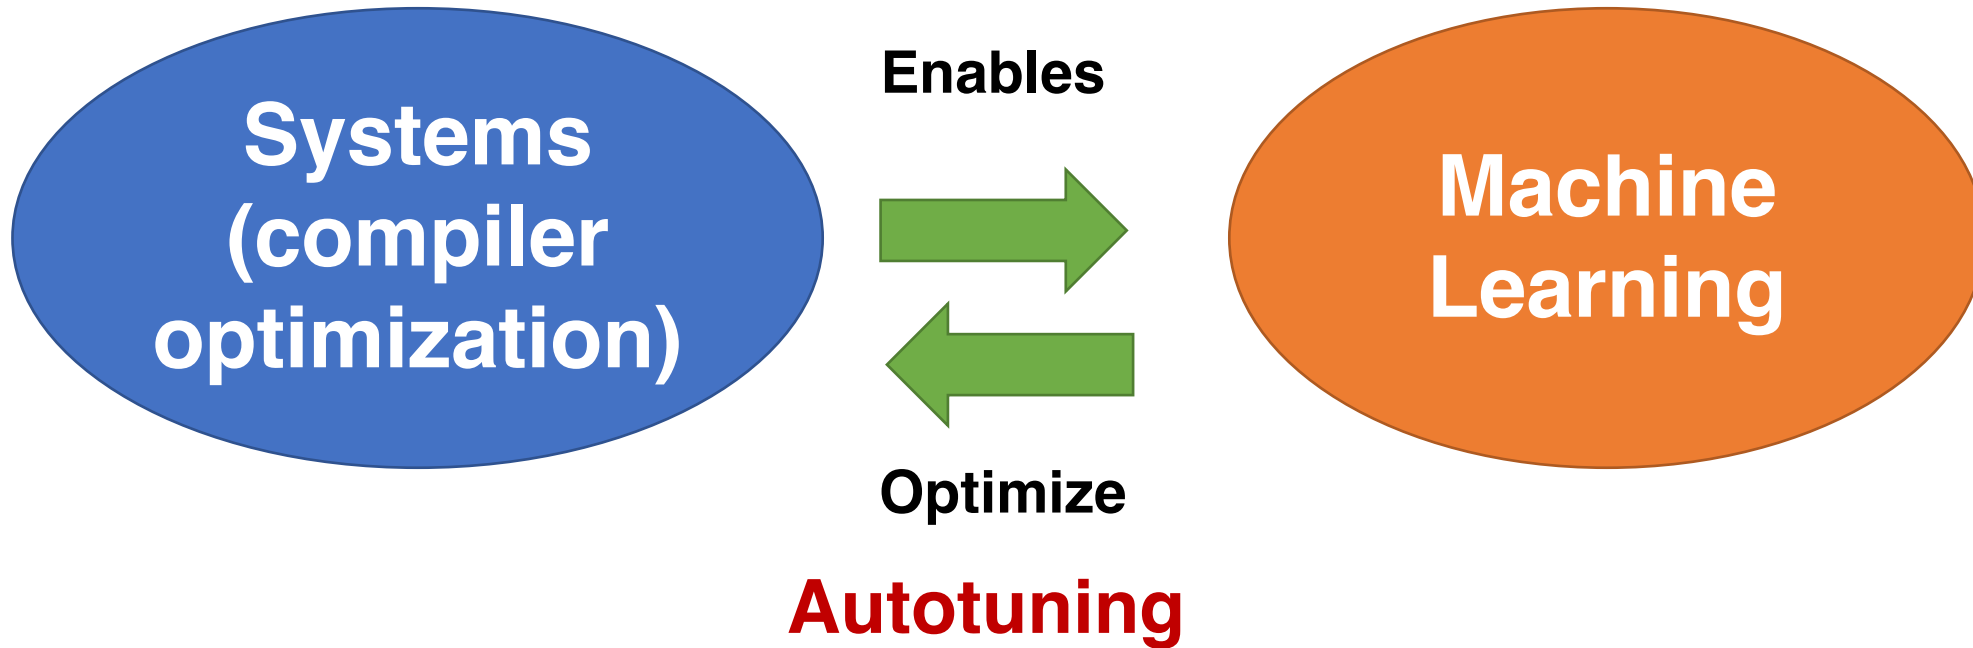
Current Learning Systems



Learning-based Learning Systems



Learning-based Learning Systems




Learning to optimize tensor programs

Learning to optimize tensor programs

- **Why do we need machine learning for systems**
- How to build intelligent systems with learning

Problem: Deep Learning Deployment





Model



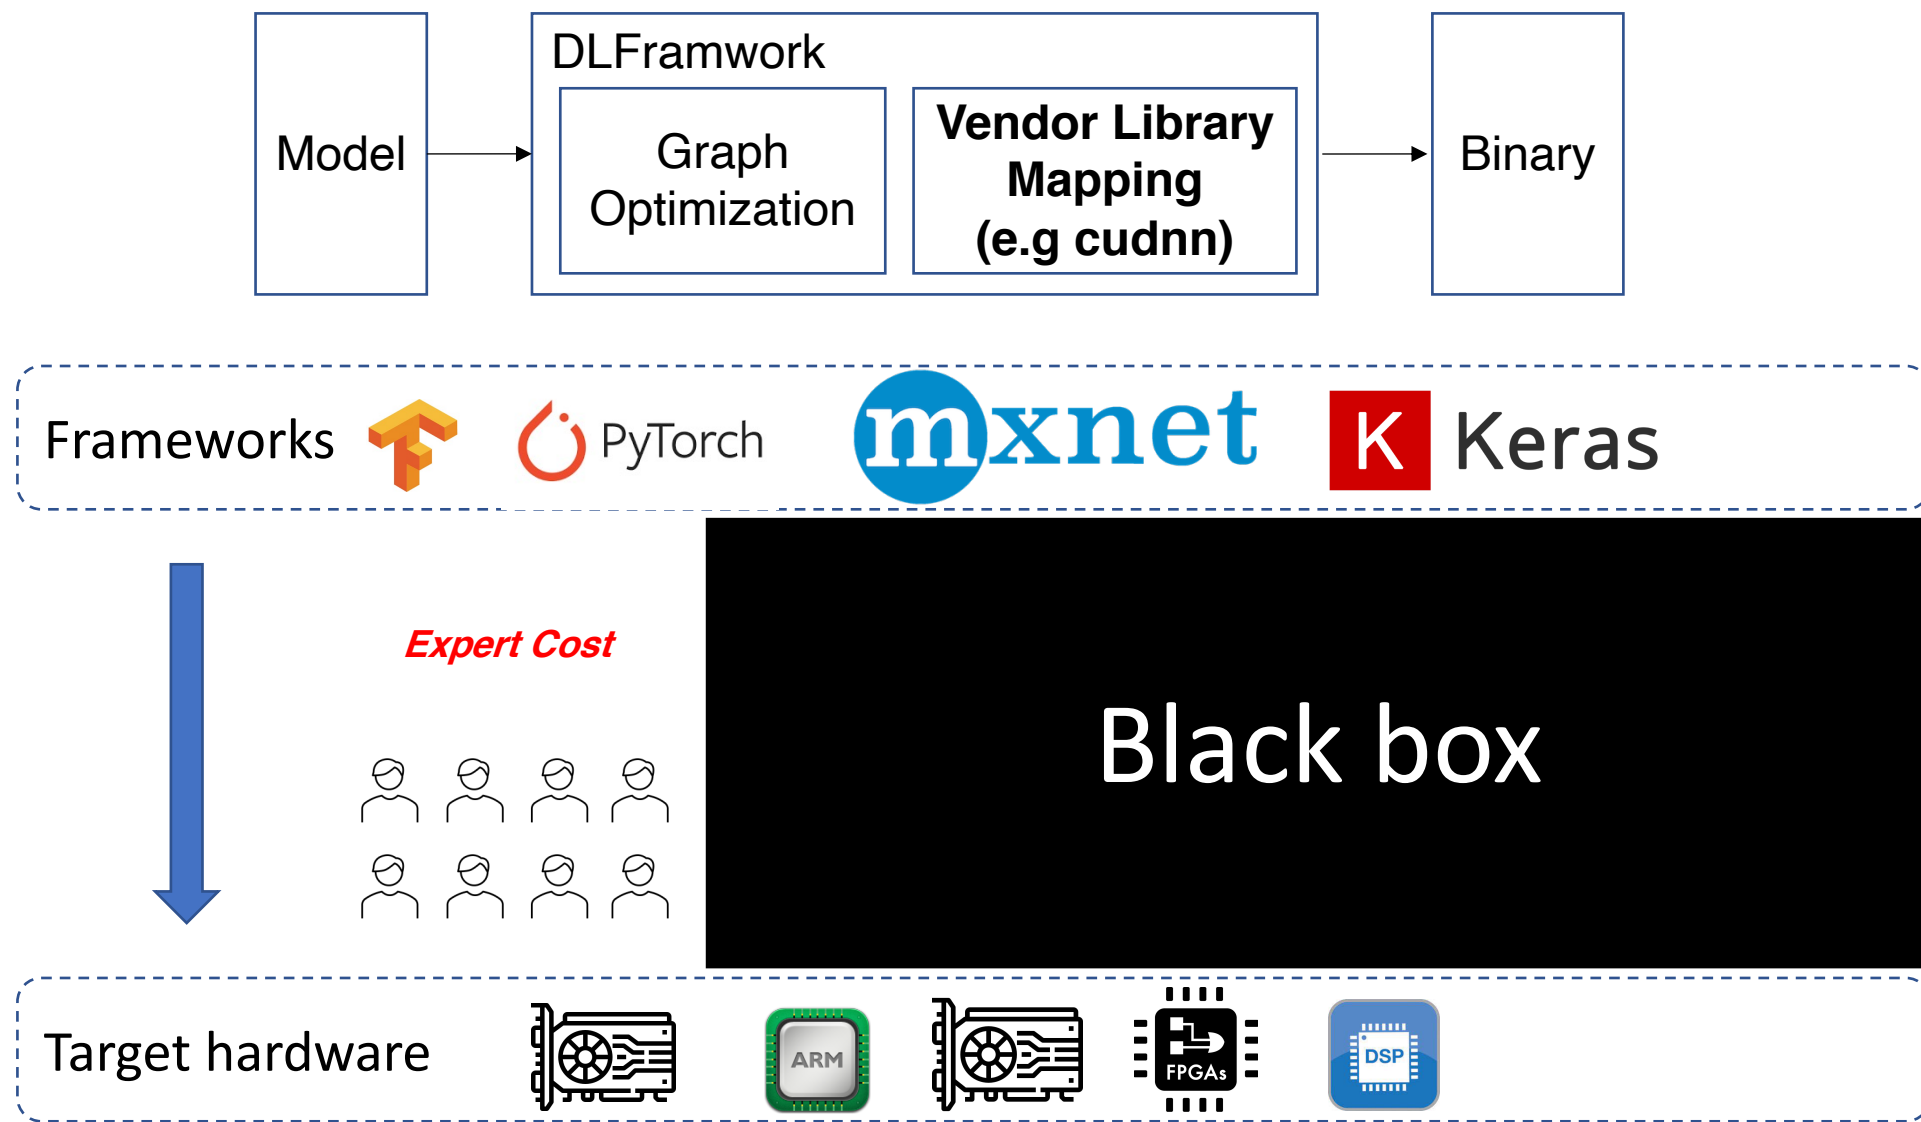
Deploy



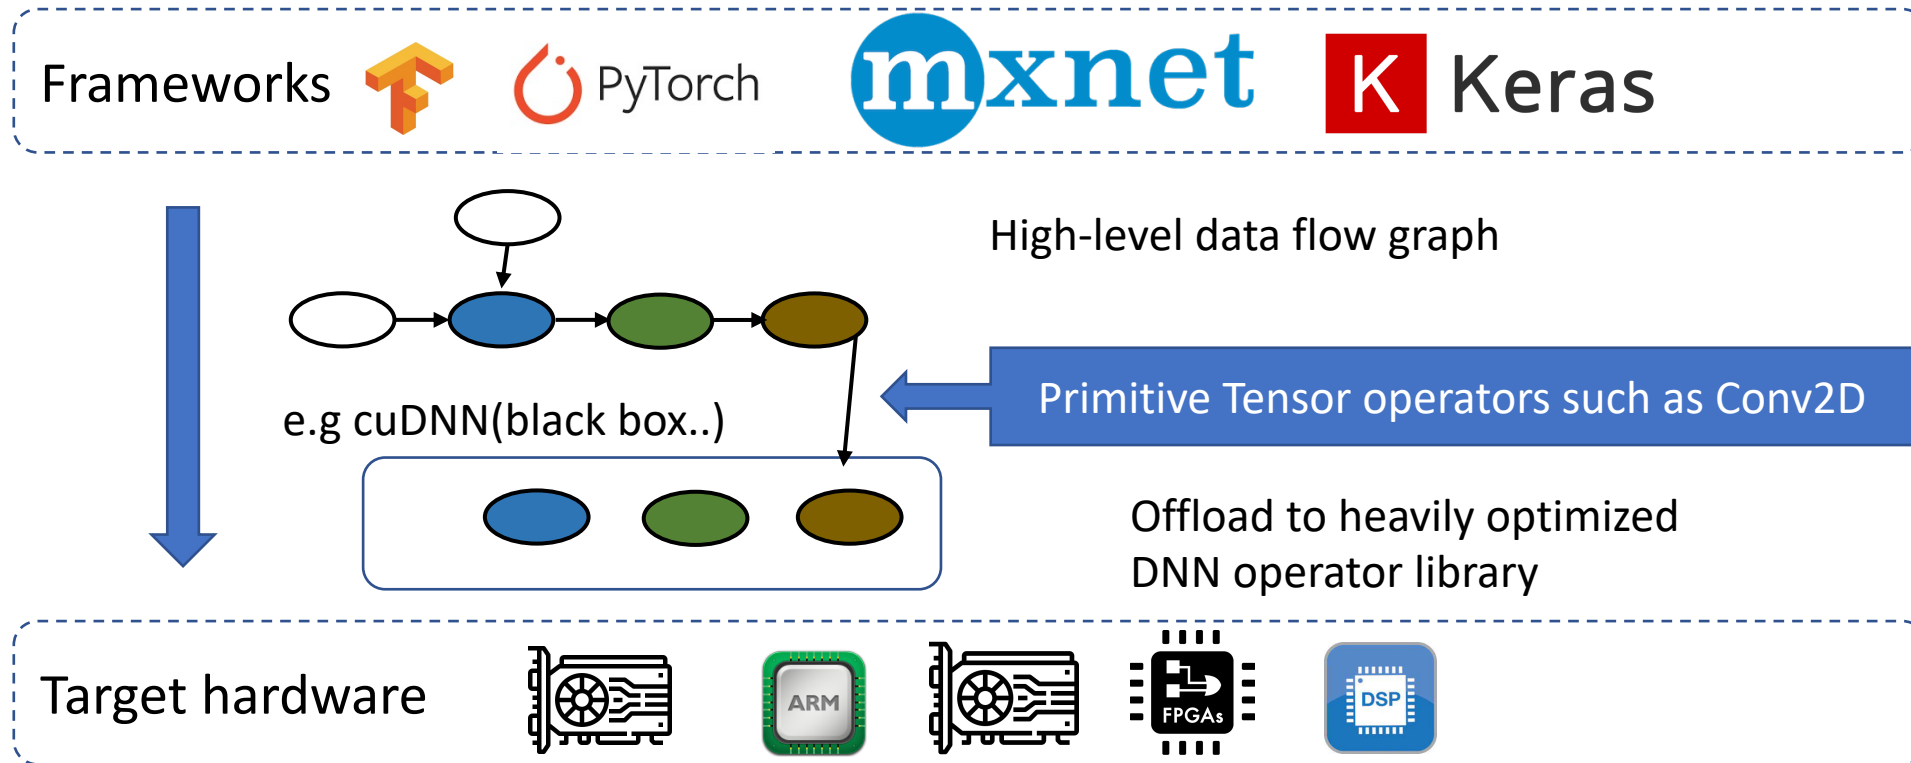
Hardware
Backends



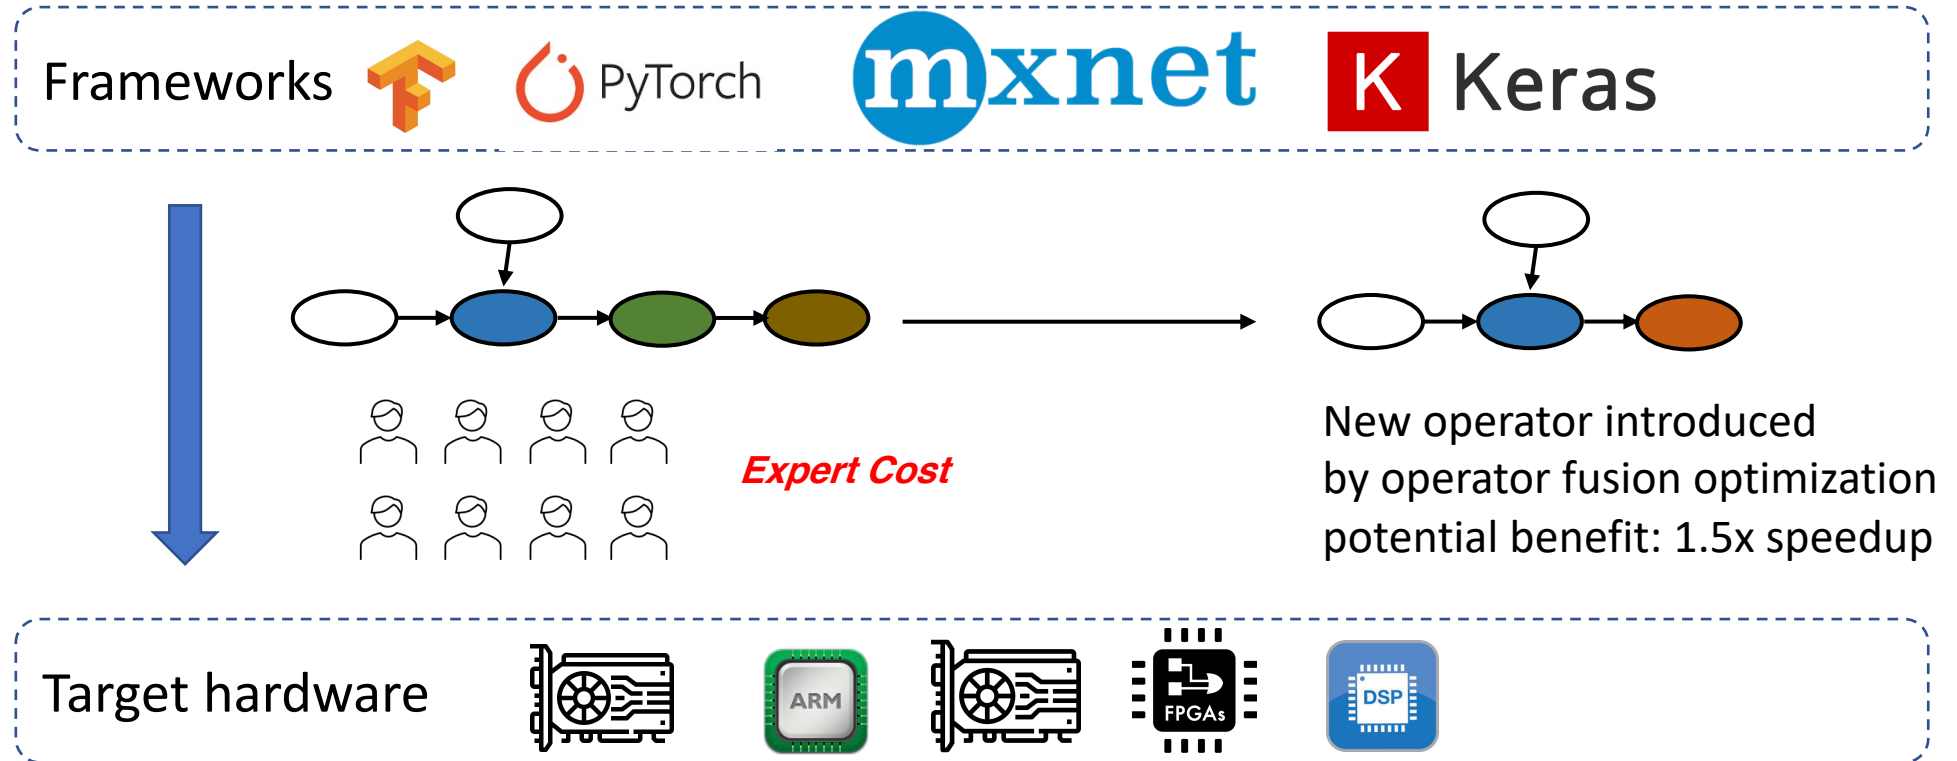
Existing Deep Learning Frameworks



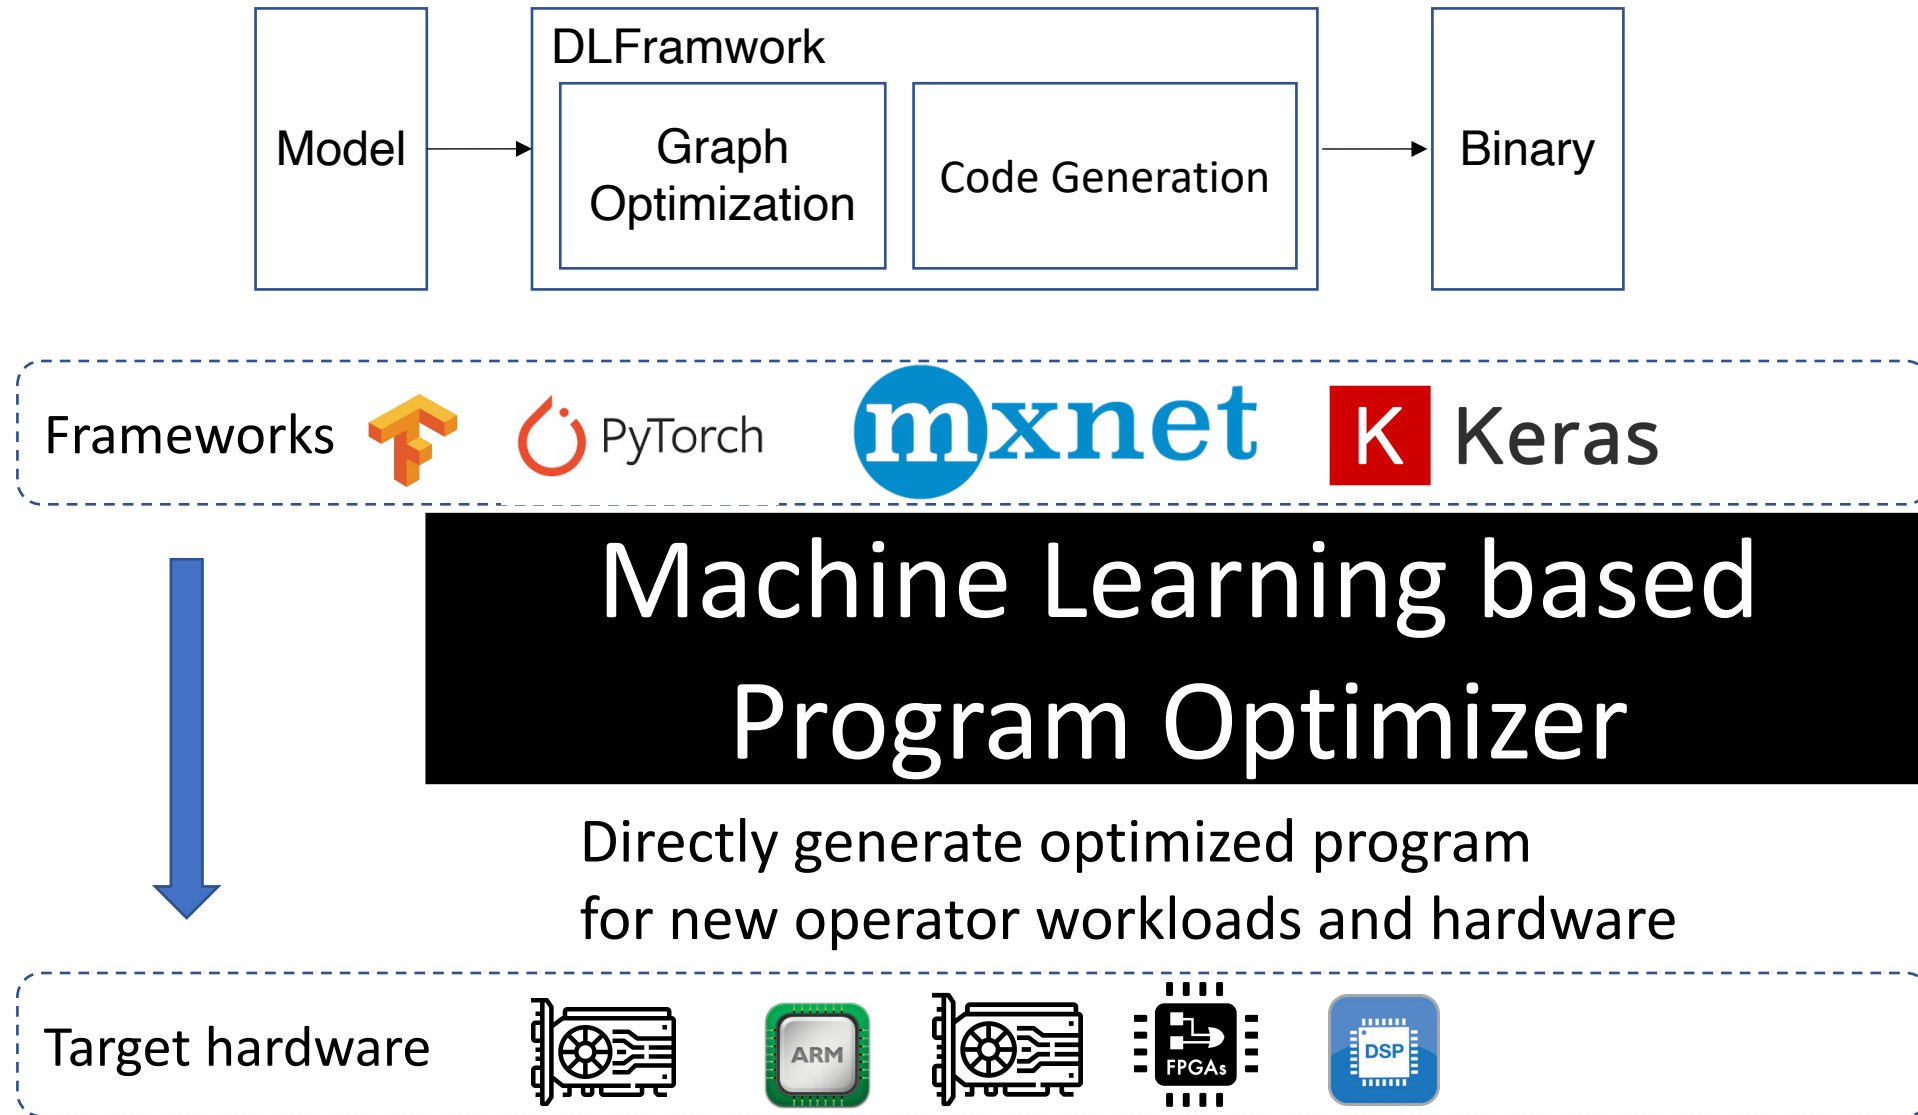
Existing Deep Learning Frameworks



Existing Deep Learning Frameworks



Learning-based Learning System



Learning to optimize tensor programs

- Why do we need machine learning for systems
- **How to build intelligent systems with learning**

Problem Setting

Tensor Expression(high level expression)

```
C = tvm.compute((m, n),  
                lambda y, x: tvm.sum(A[k, y] * B[k, x], axis=k))
```

Lowering



Loop Transformations

Thread Bindings

Cache Locality

Thread Cooperation

Tensorization

Latency Hiding

```
for y in range(1024):  
    for x in range(1024):  
        C[y][x] = 0  
        for k in range(1024):  
            C[y][x] += A[k][y] * B[k][x]
```

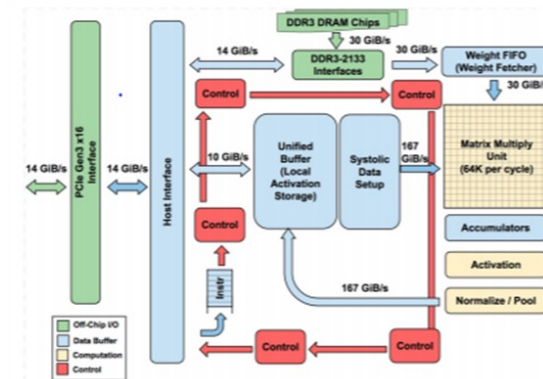
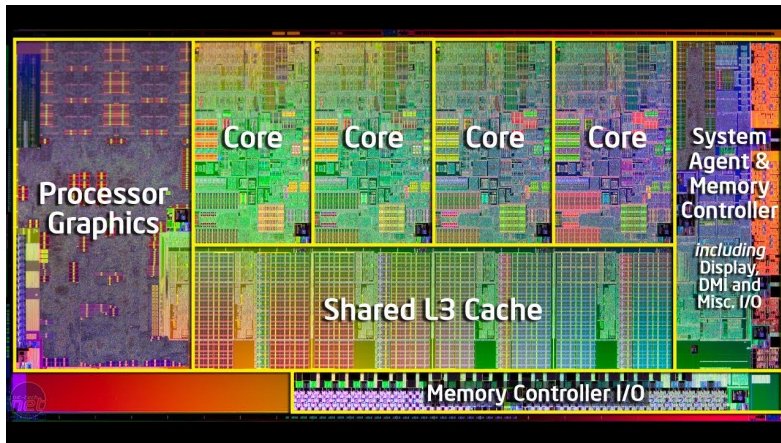
```
for yo in range(128):  
    for xo in range(128):  
        C[yo*8:yo*8+8][xo*8:xo*8+8] = 0  
        for ko in range(128):  
            for yi in range(8):  
                for xi in range(8):  
                    for ki in range(8):  
                        C[yo*8+yi][xo*8+xi] +=  
                            A[ko*8+ki][yo*8+yi] * B[ko*8+ki][xo*8+xi]
```

```
inp_buffer AL[8][8], BL[8][8]  
acc_buffer CL[8][8]  
for yo in range(128):  
    for xo in range(128):  
        vdl.a.fill_zero(CL)  
        for ko in range(128):  
            vdl.a.dma_copy2d(AL, A[ko*8:ko*8+8][yo*8:yo*8+8])  
            vdl.a.dma_copy2d(BL, B[ko*8:ko*8+8][xo*8:xo*8+8])  
            vdl.a.fused_gemm8x8_add(CL, AL, BL)  
            vdl.a.dma_copy2d(C[yo*8:yo*8+8,xo*8:xo*8+8], CL)
```


Optimization Choices in a Search Space



It is hard to consider *all hardware characteristics*.
=> **Template-based autotune**



Systolic Array on the Google TPU

Optimization Choices in a Search Space

```
@autotvm.template("tutorial/conv2d_no_batching")
def conv2d_no_batching(N, H, W, CO, CI, KH, KW, stride, padding):
    assert N == 1, "Only consider batch_size = 1 in this template"

    data = te.placeholder((N, CI, H, W), name="data")
    kernel = te.placeholder((CO, CI, KH, KW), name="kernel")
    conv = topi.nn.conv2d_nchw(data, kernel, stride, padding, dilation=1, out_dtype="float32")
    s = te.create_schedule([conv.op])
```

space definition begin


```
n, f, y, x = s[conv].op.axis
rc, ry, rx = s[conv].op.reduce_axis
```

```
cfg = autotvm.get_config()
cfg.define_split("tile_f", f, num_outputs=4)
cfg.define_split("tile_y", y, num_outputs=4)
cfg.define_split("tile_x", x, num_outputs=4)
cfg.define_split("tile_rc", rc, num_outputs=3)
cfg.define_split("tile_ry", ry, num_outputs=3)
cfg.define_split("tile_rx", rx, num_outputs=3)
cfg.define_knob("auto_unroll_max_step", [0, 512, 1500])
cfg.define_knob("unroll_explicit", [0, 1])
```

space definition end

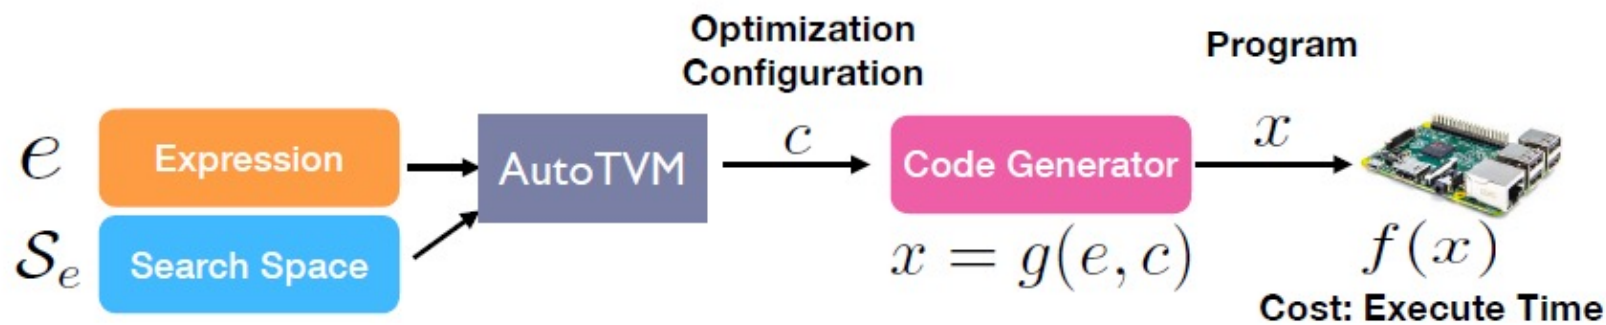
inline padding

```
pad_data = s[conv].op.input_tensors[0]
s[pad_data].compute_inline()
data, raw_data = pad_data, data
```



```
// attr [iter_var(nn.outer, )] pragma_auto_unroll_max_step = 0
// attr [iter_var(nn.outer, )] pragma_unroll_explicit = 0
for (nn.outer, 0, 1) {
    // attr [iter_var(blockIdx.z, , blockIdx.z)] thread_extent = 4
    // attr [iter_var(blockIdx.y, , blockIdx.y)] thread_extent = 3
    // attr [iter_var(blockIdx.x, , blockIdx.x)] thread_extent = 3
    // attr [iter_var(vthread, , vthread)] virtual_thread = 2
    // attr [iter_var(vthread, , vthread)] virtual_thread = 1
    // attr [iter_var(vthread, , vthread)] virtual_thread = 1
    // attr [iter_var(threadIdx.z, , threadIdx.z)] thread_extent = 1
    // attr [iter_var(threadIdx.y, , threadIdx.y)] thread_extent = 1
    // attr [iter_var(threadIdx.x, , threadIdx.x)] thread_extent = 1
    // attr [compute.local] storage_scope = "local"
    allocate compute.local[float32 * 1 * 1 * 1 * 1]
    for (rc.outer, 0, 4) {
        for (ry.outer, 0, 3) {
            for (rx.outer, 0, 3) {
                // attr [pad_temp.shared] storage_scope = "shared"
                allocate pad_temp.shared[float32 * 1 * 1 * 1 * 1]
                // attr [iter_var(threadIdx.z, , threadIdx.z)] thread_extent = 1
                // attr [iter_var(threadIdx.y, , threadIdx.y)] thread_extent = 1
                // attr [iter_var(threadIdx.x, , threadIdx.x)] thread_extent = 1
                for (ax0.ax1.fused.ax2.fused.ax3.fused.inner.inner.inner, 0, 1) {
                    pad_temp.shared[0] = placeholder[(((rc.outer*64) + (blockIdx.y*16)) + (ry.outer
                )
                // attr [placeholder.shared] storage_scope = "shared"
                allocate placeholder.shared[float32 * 2 * 1 * 1 * 1]
                // attr [iter_var(threadIdx.z, , threadIdx.z)] thread_extent = 1
                // attr [iter_var(threadIdx.y, , threadIdx.y)] thread_extent = 1
                // attr [iter_var(threadIdx.x, , threadIdx.x)] thread_extent = 1
                for (ax0.ax1.fused.ax2.fused.ax3.fused.inner.inner.inner, 0, 2) {
                    placeholder.shared[ax0.ax1.fused.ax2.fused.ax3.fused.inner.inner.inner] = placeh
                }
                for (rc.inner, 0, 1) {
                    for (ry.inner, 0, 1) {
                        for (rx.inner, 0, 1) {
                            for (nn.c, 0, 1) {
                                for (ff.c, 0, 1) {
                                    for (yy.c, 0, 1) {
                                        for (xx.c, 0, 1) {
                                            compute.local[0] = (compute.local[0] + (pad_temp.shared[0]*placeholder
                                        )
                                    }
                                }
                            }
                        }
                    }
                }
            }
        }
    }
}
```


Problem Formalization



Objective $\operatorname{argmin}_{c \in \mathcal{S}_e} f(g(e, c))$

Black-box Optimization

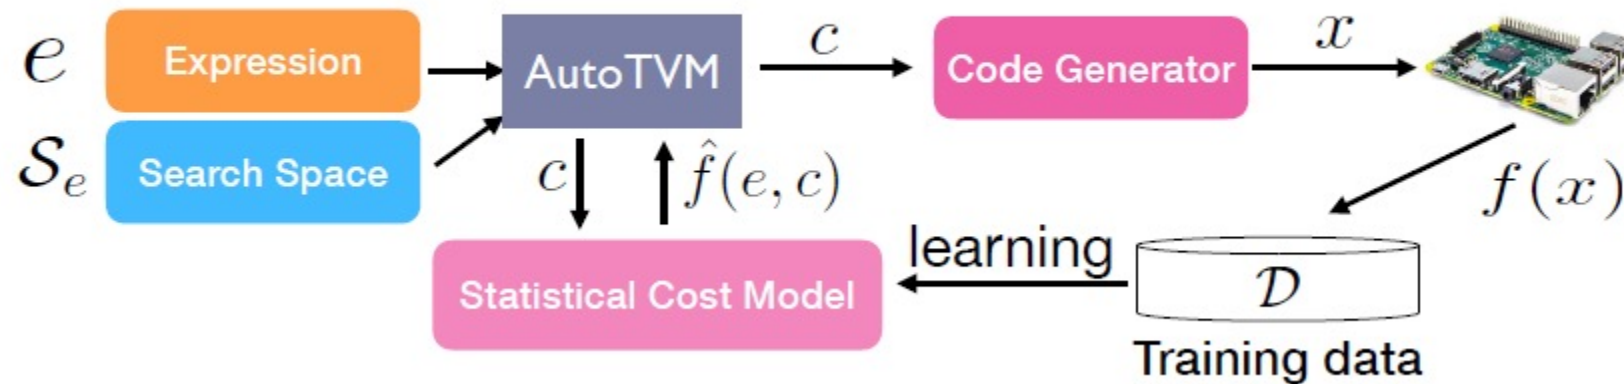
Try each configuration until we find a good one



Challenge: lots of experimental trials, each trial costs ~ 1 second

Statistical Cost Model

Use machine learning to learn a statistical cost model



Benefit: Automatically adapt to hardware type

Challenge: How to design the cost model

Loop Context Feature

```
for y in range(8):
  for x in range(8):
    C[y][x]=0
    for k in range(8):
      C[y][x]+=A[k][y]*B[k][x]
```

(a) Low level AST

	touched memory			outer loop length
	C	A	B	
y	64	64	64	1
x	8	8	64	8
k	1	8	8	64

(b) Loop context vectors

Flatten as
feature vector

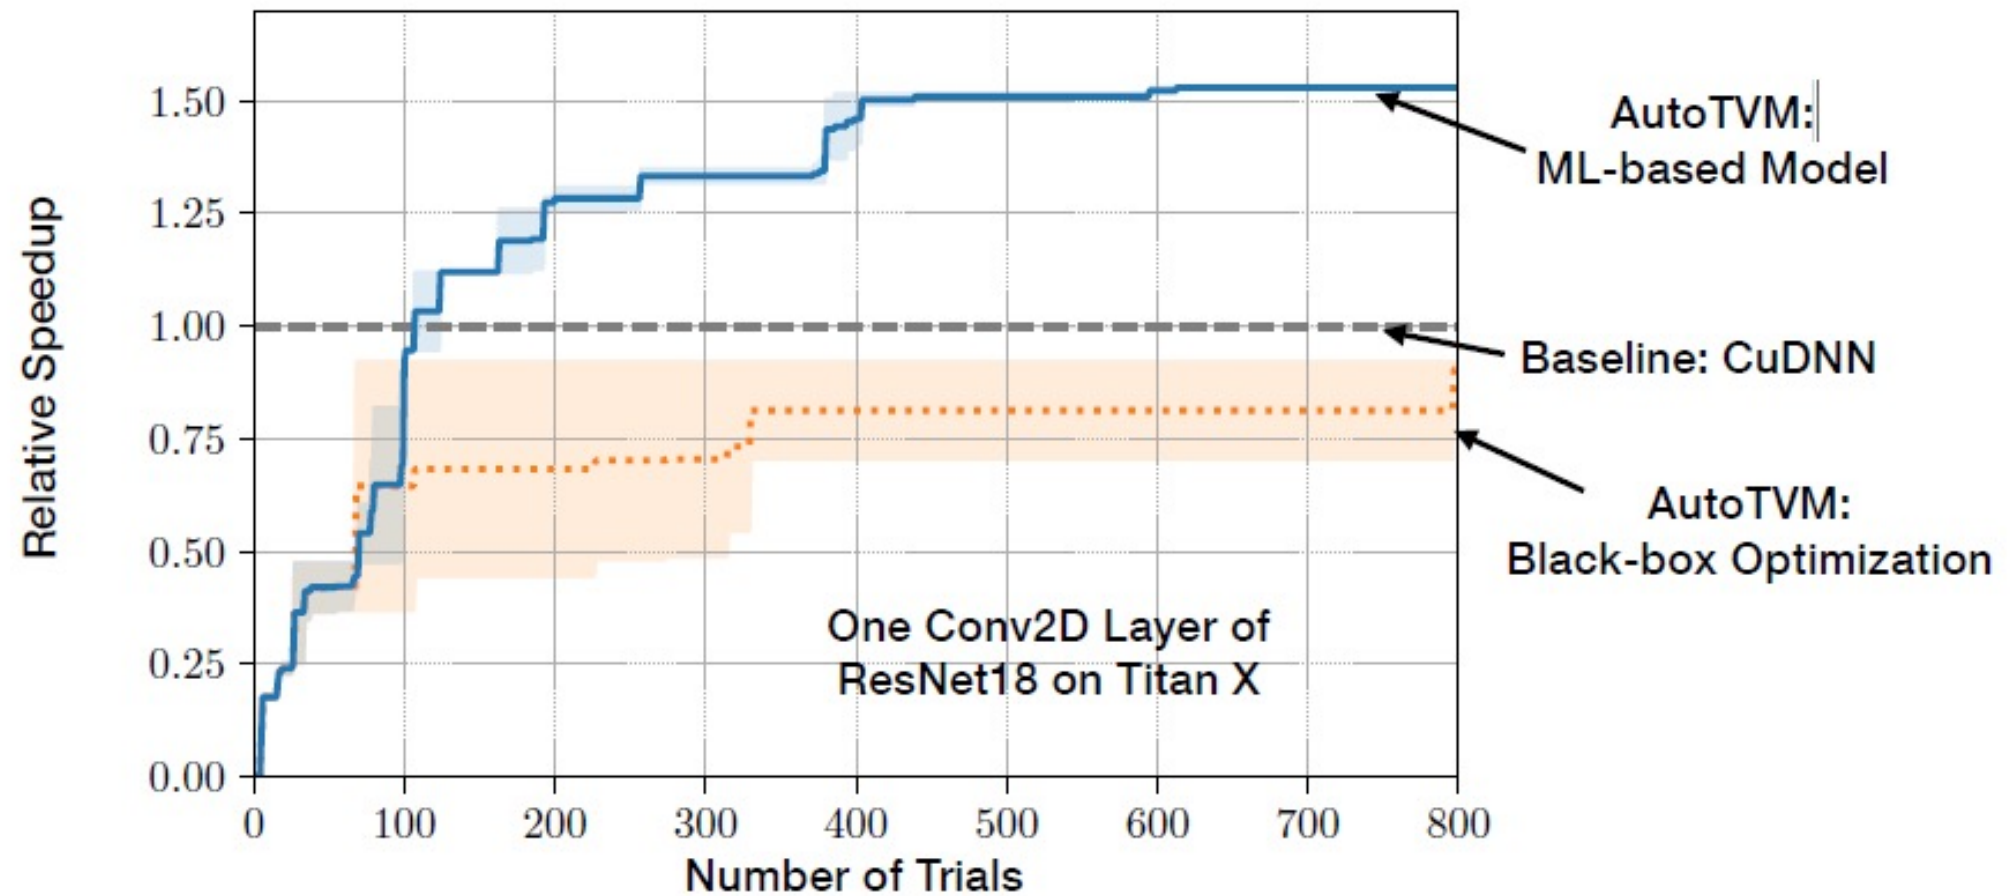


Feature Vector

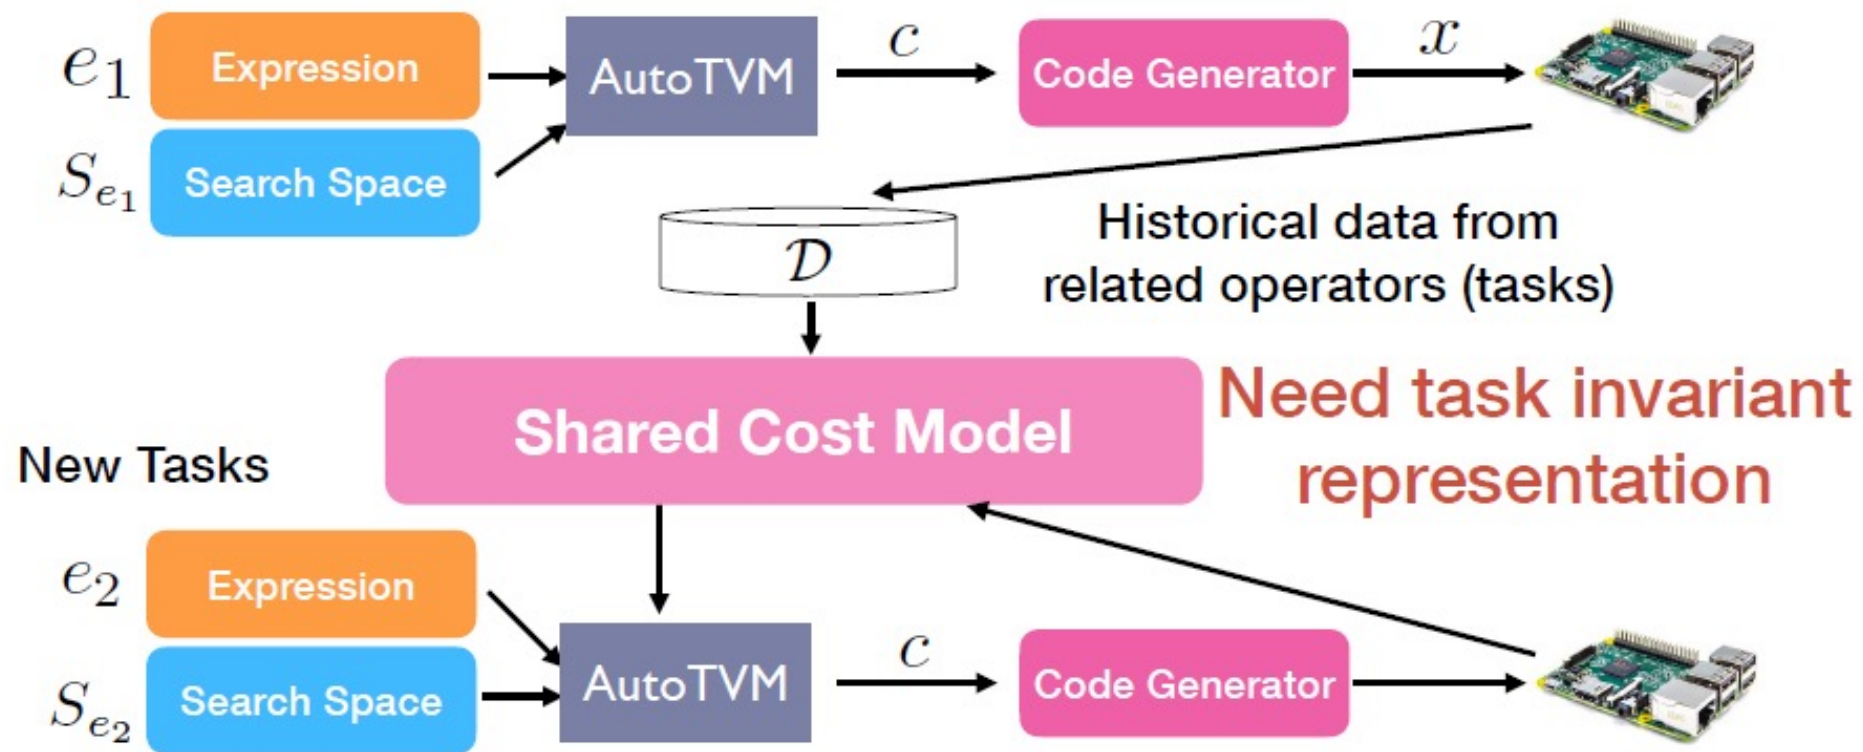
Feature Name		Description
length		The length of this loop
annotation		One-hot annotation of this loop (can be vectorize, unrolled, paralleled, ...)
top-down		The product of the lengths of outer loops
bottom-up		The product of the lengths of inner loops
access pattern (for every buffer)	touch count	The number of touched elements
	reuse ratio	Reuse ratio of this buffer (= bottom-up / touch count)
	stride	Coefficient of this loop variable in the index expression

Table 2: Listing of loop context feature

Effectiveness of ML based Model



Transferable Cost Model



Impact of Transfer Learning




Figure 8: Impact of transfer learning. Transfer-based models quickly found better solutions.

Mxnet : v1.1
 TF-GPU:v1.7
 TFLite:7558b085
 ARM Compute
 Library :v18.03




Figure 11: End-to-end performance across back-ends. ²AutoTVM outperforms the baseline methods.

Ansor: Generating High- Performance Tensor Programs for Deep Learning

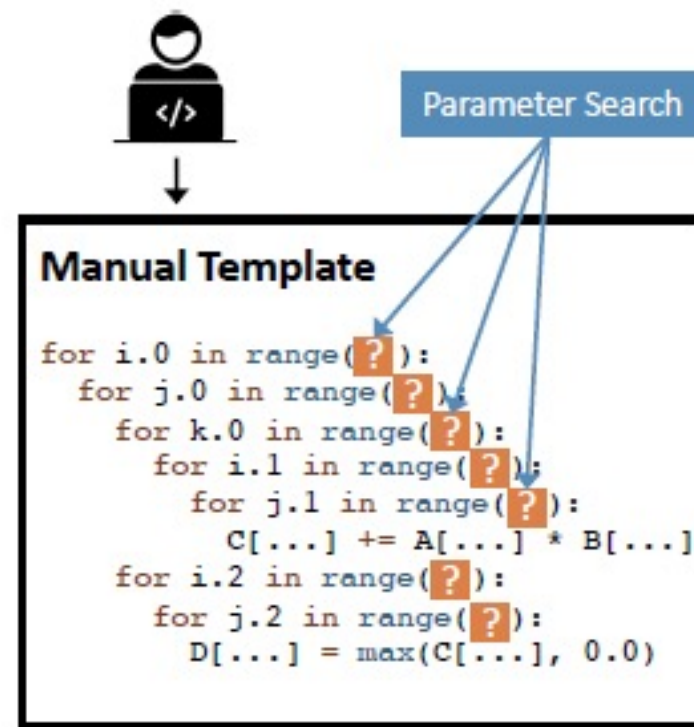
TVM's Approach

AutoTVM: Template-guided search

Use **templates** to define the search space for every operator

Drawbacks

- Not fully-automated -> Requires huge manual effort(15K lines of code)
- Limited search space -> Does not achieve optimal performance



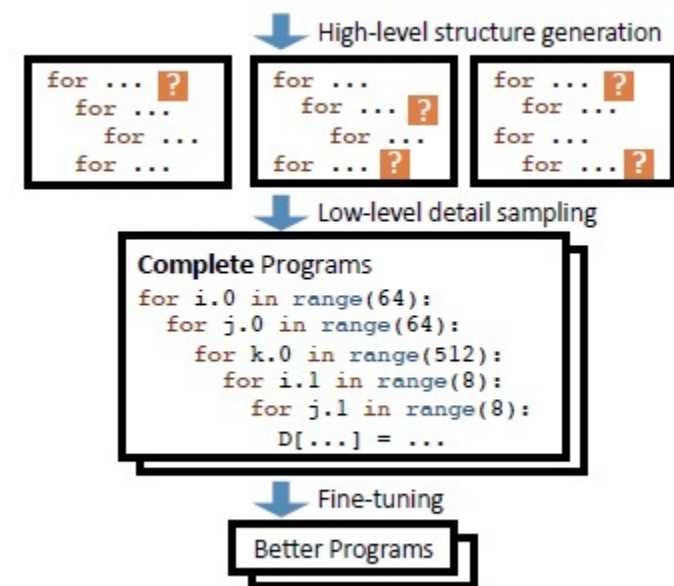
Challenges and ansor's approach

C1: How to build a large search space automatically?

- Use a hierarchical search space

C2: How to search efficiently?

- Sample complete programs and fine-tune them




Challenges and ansor's approach

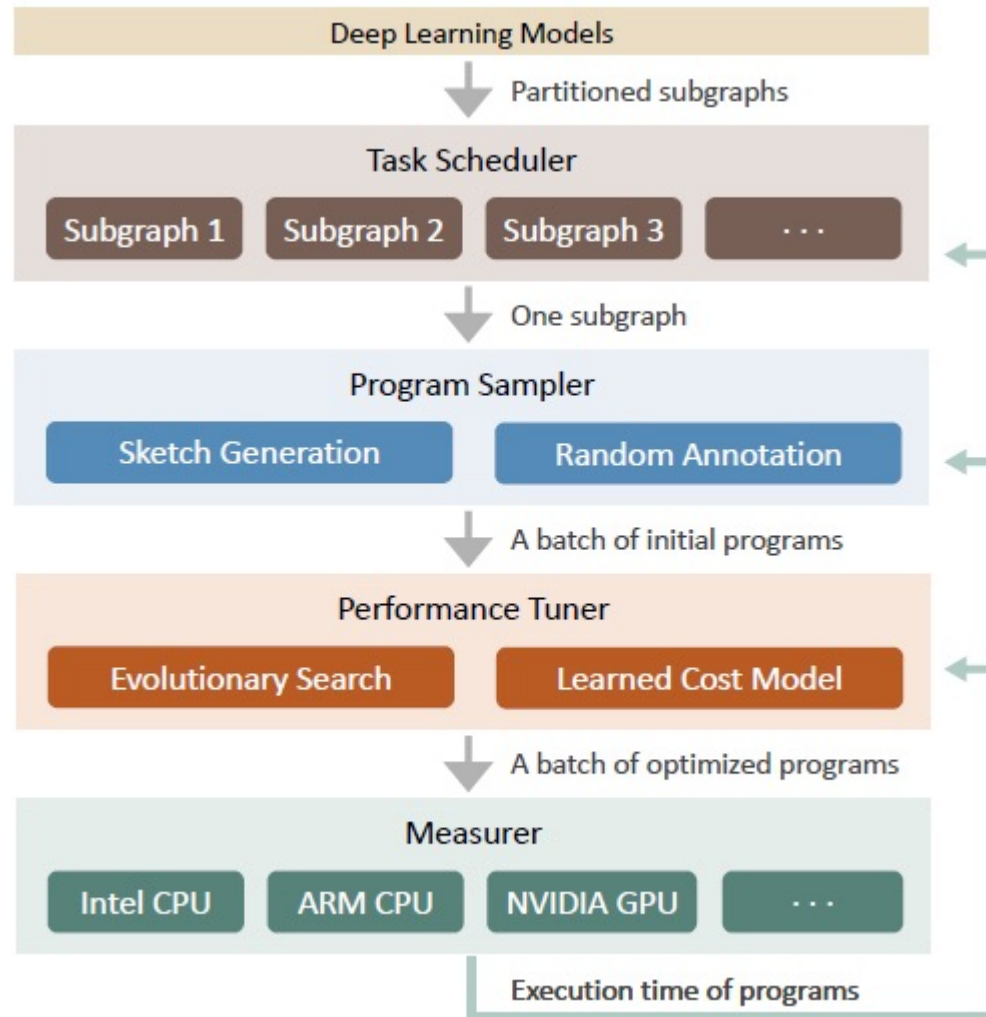
Need to generate programs for all layers -> A lot of search tasks

C3: How to allocate resource for many search tasks?

- Utilize a task scheduler to prioritize important tasks



Ansor Overview



Program Sampling

- **Goal:** automatically construct a large search space and uniformly sample from the space
- **Approach**
 - Two-level hierarchical search space: **Sketch** + **Annotation**
 - **Sketch:** a few good high-level structures
 - **Annotation:** billions of low-level details

Sampling Process



Sketch Generation Examples

Example Input 1:

* **The mathematical expression:**

$$C[i, j] = \sum_k A[i, k] \times B[k, j]$$


$$D[i, j] = \max(C[i, j], 0.0)$$

where $0 \leq i, j, k < 512$

* **The corresponding naïve program:**

```
for i in range(512):
    for j in range(512):
        for k in range(512):
            C[i, j] += A[i, k] * B[k, j]
for i in range(512):
    for j in range(512):
        D[i, j] = max(C[i, j], 0.0)
```

* **The corresponding DAG:**



Derivation:

$$\text{Input 1} \rightarrow \sigma(S_0, i=4) \xrightarrow{\text{Rule 1}} \sigma(S_1, i=3) \xrightarrow{\text{Rule 4}}$$

$$\sigma(S_2, i=2) \xrightarrow{\text{Rule 1}} \sigma(S_3, i=1) \xrightarrow{\text{Rule 1}} \text{Sketch 1}$$

Generated sketch 1

```
for i.0 in range(TILE_I0):
    for j.0 in range(TILE_J0):
        for i.1 in range(TILE_I1):
            for j.1 in range(TILE_J1):
                for k.0 in range(TILE_K0):
                    for i.2 in range(TILE_I2):
                        for j.2 in range(TILE_J2):
                            for k.1 in range(TILE_I1):
                                for i.3 in range(TILE_I3):
                                    for j.3 in range(TILE_J3):
                                        C[...] += A[...] * B[...]
for i.4 in range(TILE_I2 * TILE_I3):
    for j.4 in range(TILE_J2 * TILE_J3):
        D[...] = max(C[...], 0.0)
```

“SSRSRSS” multi-level tiling + fusion

Sketch Generation Examples

Example Input 1:

*** The mathematical expression:**

$$C[i, j] = \sum_k A[i, k] \times B[k, j]$$

$$D[i, j] = \max(C[i, j], 0.0)$$

where $0 \leq i, j, k < 512$

*** The corresponding naïve program:**

```
for i in range(512):
    for j in range(512):
        for k in range(512):
            C[i, j] += A[i, k] * B[k, j]
for i in range(512):
    for j in range(512):
```



Derivation:

$$\text{Input 1} \rightarrow \sigma(S_0, i=4) \xrightarrow{\text{Rule 1}} \sigma(S_1, i=3) \xrightarrow{\text{Rule 4}}$$

$$\sigma(S_2, i=2) \xrightarrow{\text{Rule 1}} \sigma(S_3, i=1) \xrightarrow{\text{Rule 1}} \text{Sketch 1}$$

Generated sketch 1

```
for i.0 in range(TILE_I0):
    for j.0 in range(TILE_J0):
        for i.1 in range(TILE_I1):
            for j.1 in range(TILE_J1):
                for k.0 in range(TILE_K0):
                    for i.2 in range(TILE_I2):
                        for j.2 in range(TILE_J2):
                            for k.1 in range(TILE_I1):
                                for i.3 in range(TILE_I3):
                                    for j.3 in range(TILE_J3):
                                        C[...] += A[...] * B[...]
or i.4 in range(TILE_I2 * TILE_I3):
    for j.4 in range(TILE_J2 * TILE_J3):
        D[...] = max(C[...], 0.0)
```

"SSRSRSS" multi-level tiling + fusion

No	Rule Name	Condition
1	Skip	$\neg \text{IsStrictInlinable}(S, i)$
2	Always Inline	$\text{IsStrictInlinable}(S, i)$
3	Multi-level Tiling	$\text{HasDataReuse}(S, i)$
4	Multi-level Tiling with Fusion	$\text{HasDataReuse}(S, i) \wedge \text{HasFusibleConsumer}(S, i)$
5	Add Cache Stage	$\text{HasDataReuse}(S, i) \wedge \neg \text{HasFusibleConsumer}(S, i)$
6	Reduction Factorization	$\text{HasMoreReductionParallel}(S, i)$
...	User Defined Rule	...

Random Annotation Examples

- Parallelize some outer loop
- Vectorize some inner loop
- unroll few inner loop
- randomly fill tile size

Generated sketch 1

```
for i.0 in range(TILE_I0):
  for j.0 in range(TILE_J0):
    for i.1 in range(TILE_I1):
      for j.1 in range(TILE_J1):
        for k.0 in range(TILE_K0):
          for i.2 in range(TILE_I2):
            for j.2 in range(TILE_J2):
              for k.1 in range(TILE_I1):
                for i.3 in range(TILE_I3):
                  for j.3 in range(TILE_J3):
                    C[...] += A[...] * B[...]
          for i.4 in range(TILE_I2 * TILE_I3):
            for j.4 in range(TILE_J2 * TILE_J3):
              D[...] = max(C[...], 0.0)
```

Sampled program 1

```
parallel i.0@j.0@i.1@j.1 in range(256):|
  for k.0 in range(32):
    for i.2 in range(16):
      unroll k.1 in range(16):
        unroll i.3 in range(4):
          vectorize j.3 in range(16):
            C[...] += A[...] * B[...]
    for i.4 in range(64):
      vectorize j.4 in range(16):
        D[...] = max(C[...], 0.0)
```

Sampled program 2

```
parallel i.2 in range(16):
  for j.2 in range(128):
    for k.1 in range(512):
      for i.3 in range(32):
        vectorize j.3 in range(4):
          C[...] += A[...] * B[...]
parallel i.4 in range(512):
  for j.4 in range(512):
    D[...] = max(C[...], 0.0)
```

Learned Cost Model

Predict the score of each non-loop innermost statement

Example:


Statement B:	<pre>for i in range(10): for j in range(10): B[i][j] = A[i] * 2</pre>
Statement C:	<pre>for i in range(10): C[i] = B[i][i] - 3</pre>

Cost = Cost of Statement B + Cost of Statement C


- Extract features for every non-loop innermost statement:
 - used cache lines, used memory, reuse distance, arithmetic intensity, ...
- Train on-the-fly with measured programs (typically less than 30,000)

Task Scheduler

- There are many **subgraphs** (search tasks) in a network
 - Example: ResNet-50 has 29 unique subgraphs after partition
- **Existing systems:** sequential optimization with a fixed allocation



- **Our task scheduler:** slice the time and prioritize important subgraphs



- Predict each task's impact on the end-to-end objective function
 - Using optimistic guess and similarity between tasks

Single Operator

Platform:

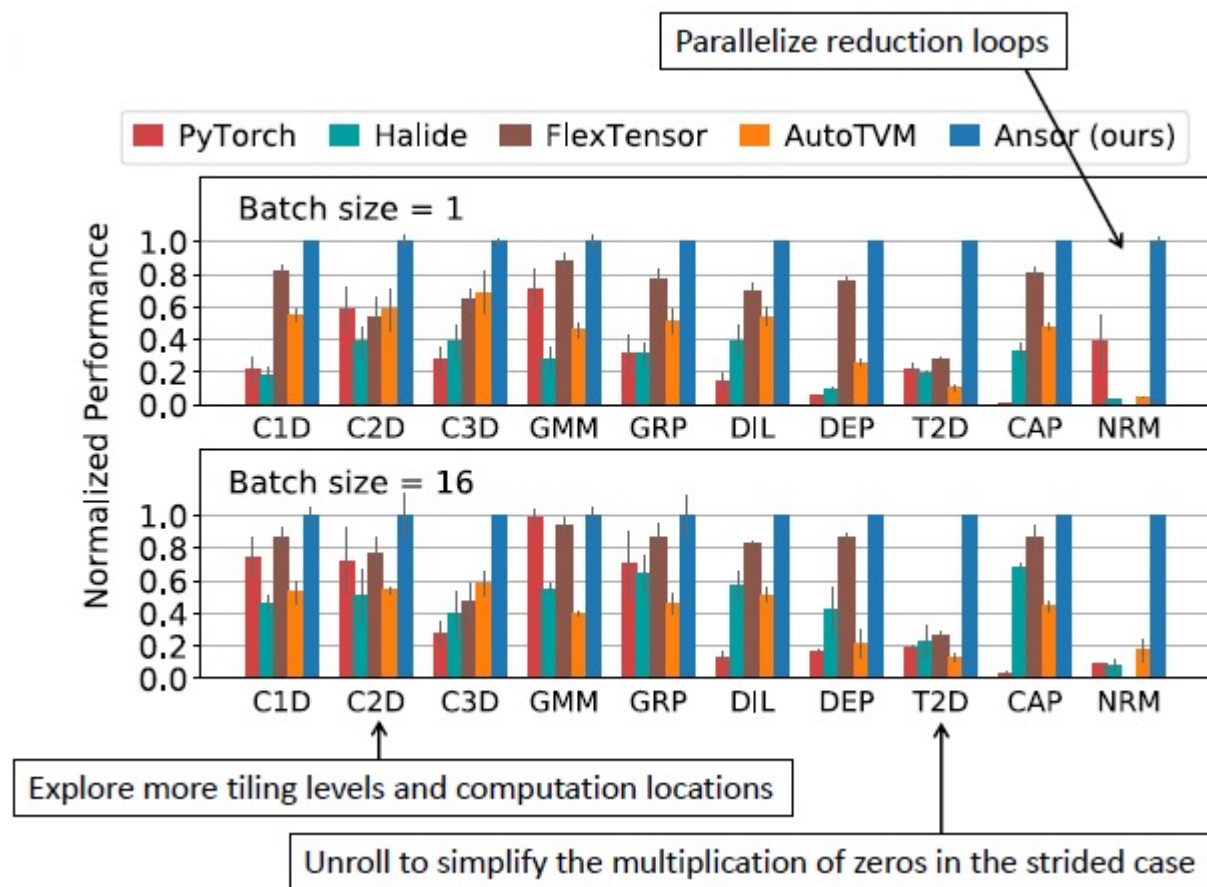
Intel-Platinum 8124M (18 cores)

Operators:

conv1d (C1D), conv2d (C2D),
conv3d (C3D), matmul (GMM)
group conv2d (GRP),
dilated conv2d (DIL)
depthwise conv2d (DEP),
conv2d transpose (T2D),
capsule conv2d (CAP),
matrix 2-norm (NRM)

Analysis:

For most test cases, the best programs found by Ansoor are outside the search space of existing search-based frameworks.



Subgraph

Platforms:

"@C" for Intel CPU (8124M)

"@G" for NVIDIA (V100)

Subgraphs:

ConvLayer = conv2d + bn + relu

TBS = transpose + batch_matmul
+ softmax

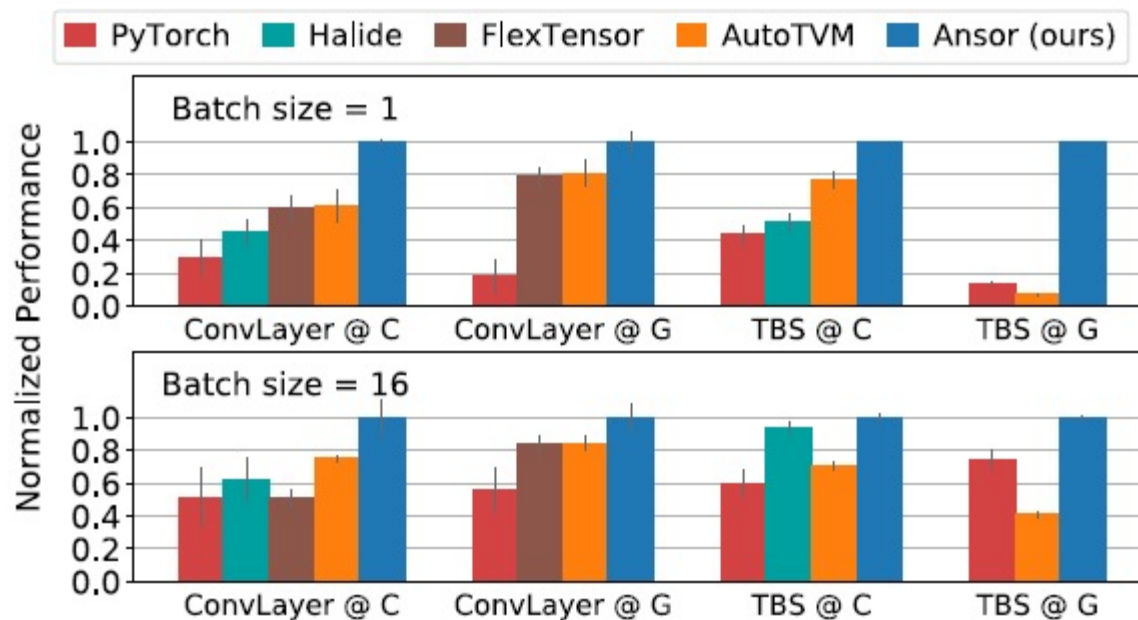
Library Version

PyTorch (v1.5 with torch script)

TensorFlow (v2.0 with graph mode)

TensorRT (v6.0 with TensorFlow integration)

TensorFlow Lite (V2.0)



Network

Platforms:

Intel CPU (8124M)

NVIDIA GPU (V100)

ARM CPU (A53)

Networks:

ResNet-50, Mobilenet V2,

3D-ResNet, DCGAN, BERT

Library Version

PyTorch (v1.5 with torch script)


TensorFlow (v2.0 with graph mode)

TensorRT (v6.0 with TensorFlow integration)

TensorFlow Lite (V2.0)

Analysis

- Ansor performs best or equally the best in all test cases with up to **3.8x** speedup



Network

Platforms:

Intel CPU (8124M)

NVIDIA GPU (V100)

ARM CPU (A53)

Networks:

ResNet-50, Mobilenet V2,

3D-ResNet, DCGAN, BERT

Library Version

PyTorch (v1.5 with torch script)


TensorFlow (v2.0 with graph mode)

TensorRT (v6.0 with TensorFlow integration)

TensorFlow Lite (V2.0)

Analysis

- Ansor performs best or equally the best in all test cases with up to **3.8x** speedup



Network

Platforms:

Intel CPU (8124M)

NVIDIA GPU (V100)

ARM CPU (A53)

Networks:

ResNet-50, Mobilenet V2,

3D-ResNet, DCGAN, BERT

Library Version

PyTorch (v1.5 with torch script)


TensorFlow (v2.0 with graph mode)

TensorRT (v6.0 with TensorFlow integration)


TensorFlow Lite (V2.0)

Analysis

- Ansor performs best or equally the best in all test cases with up to **3.8x** speedup
- Ansor delivers portable performance



Ablation Study



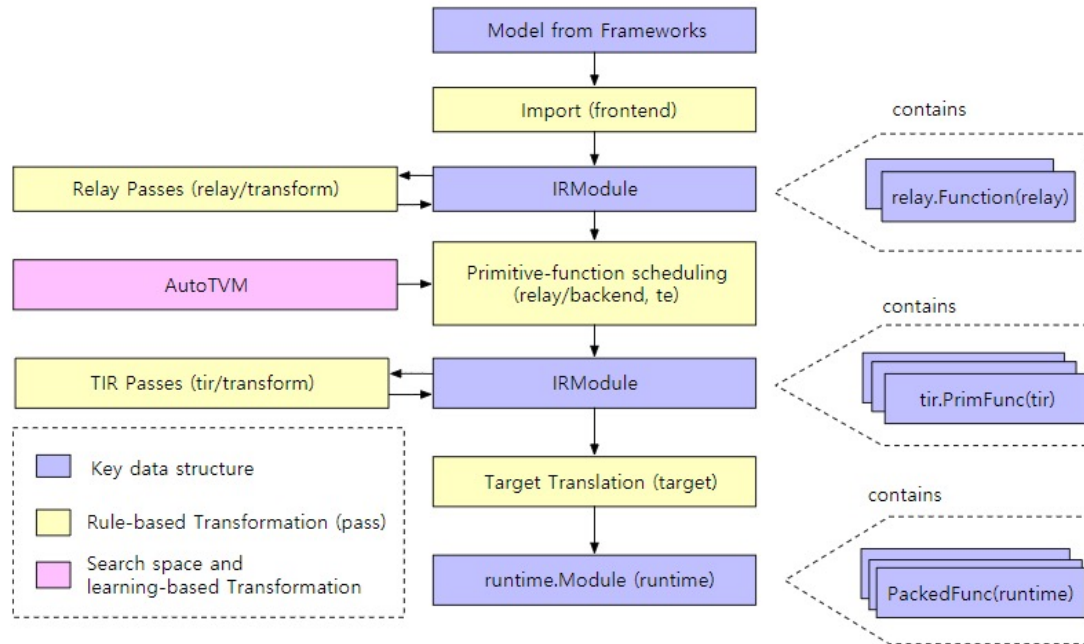
Analysis

- The most important factor is the search space
- Fine-tuning improves the search results significantly
- Task scheduler accelerates the search
- Match the performance of AutoTVM with 10x less search time

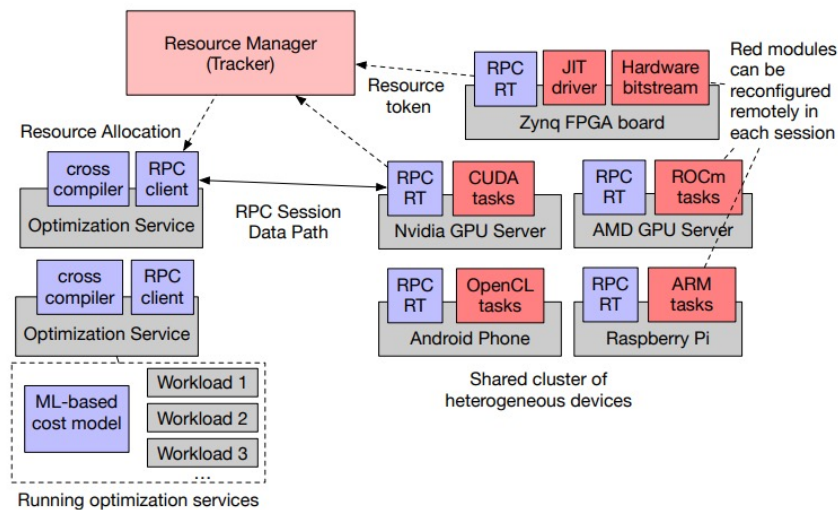
Use case

- AutoTVM
 - Well known parameter space
 - Special operation(quantization,Winograd) and hardware(tensor-core)
- AutoScheduler
 - The effect of the parameter is unclear
 - Need more potential performance gain
 - when have a problem creating a sufficient template

Use case



- AutoTVM template
- Rule based transformation
- Runtime



Use case

```
# extract workloads from relay program
print("Extract tasks...")
mod, params, input_shape, _ = get_network(network, batch_size=1)
tasks = autotvm.task.extract_from_program(
    mod["main"],
    target=target,
    params=params,
    ops=(relay.op.get("nn.conv2d"),),
)
```

Use case

```
#### DEVICE CONFIG ####
target = tvm.target.cuda()

#### TUNING OPTION ####
network = "resnet-18"
log_file = "%s.log" % network
dtype = "float32"

tuning_option = {
    "log_filename": log_file,
    "tuner": "xgb",
    "n_trial": 2000,
    "early_stopping": 600,
    "measure_option": autotvm.measure_option(
        builder=autotvm.LocalBuilder(timeout=10),
        runner=autotvm.LocalRunner(number=20, repeat=3, timeout=4, min_repeat_ms=150),
    ),
}
```

AutoTVM

```
print("Begin tuning...")
measure_ctx = auto_scheduler.LocalRPCMeasureContext(repeat=1, min_repeat_ms=300, timeout=10)

tuner = auto_scheduler.TaskScheduler(tasks, task_weights)
tune_option = auto_scheduler.TuningOptions(
    num_measure_trials=200, # change this to 20000 to achieve the best performance
    runner=measure_ctx.runner,
    measure_callbacks=[auto_scheduler.RecordToFile(log_file)],
)

tuner.tune(tune_option)
```

Ansor

Tips

- Search space design
- # of trial
- Kinds of exploration algorithm
- Cost model
- Hyper-parameter

Thanks

References

- Learning to optimize tensor programs [paper](#)
- Learning to optimize tensor programs [slide](#)
- [Dive into Deep Learning Compiler](#)
- Ansor [slide](#)
- Ansor [paper](#)
- Tvm [docs](#)