

# **DNNFusion:** Accelerating Deep Neural Networks Execution with Advanced Operator Fusion

(PLDI 2021)

Wei Niu<sup>1</sup>, Jiexiong Guan<sup>1</sup>, Yanzhi Wang<sup>2</sup>,  
Gagan Agrawal<sup>3</sup>, Bin Ren<sup>1</sup>

<sup>1</sup>William & Mary, <sup>2</sup>Northeastern University, <sup>3</sup>Augusta University

Presenter: Taehee Jeong

# Optimizing Deep Neural Networks

Deep networks need high memory and computation requirements

- **Operator fusion** is a key optimization technique in many frameworks (TF, TVM, MNN)

Existing approaches adapt **too restrictive fusion strategies** (i.e. few hand-coded pattern matching rules)

- Not able to cover **diversity** of operators and layer connections
- *Polyhedral-based* loop fusion techniques: focused on **affine-loop** optimizations, won't be able to capture some operation combinations

# Contributions

- Graph rewriting based on **mathematical property** of operations
- Operator fusion plan generation from high-level operator abstractions (with **mapping types** and mathematical properties)
- Optimized code generation for **fused** blocks

⇒ Up to **8.8x** more loop fusions and **9.3x** speedup compared to existing frameworks

# Deeper models have lesser throughput

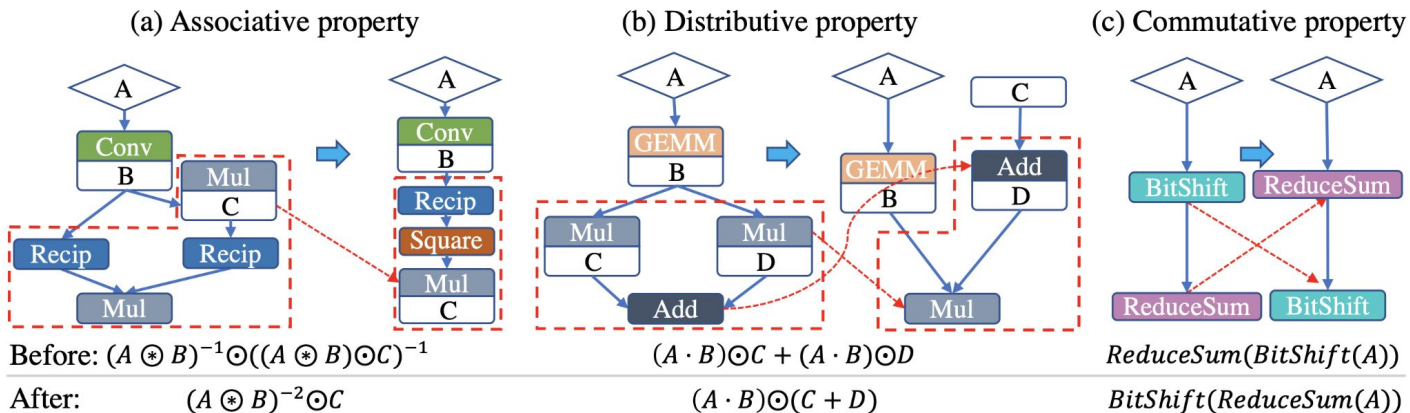
Model	#Total layer	IR size	#FLOPS	Speed (FLOPs/S)
VGG-16 [62]	51	161M	31.0B	320G
YOLO-V4 [7]	398	329M	34.6B	135G
DistilBERT [60]	457	540M	35.3B	78G
MobileBERT [65]	2,387	744M	17.6B	44G
GPT-2 [55]	2,533	1,389M	69.1B	62G

- More layer: more intermediate results
  - Memory/cache pressure ↑
  - ex) **reshape**, **squeeze** in TFLite are just **memcpy**!
- Insufficient amount of computations per layer
  - Utilization ↓ (esp. for GPU)

⇒ Operator fusion can pack more computations per (fused) layer

# Mathematical property based graph rewriting

- Mathematical properties: associative, distributive, commutative
- Identifies set of rewrite rules with **operator fusion in mind**
  - Results in fewer layers in the final fused graph
- Focus on 1-to-1 mappings and reduction operators (many-to-many)
- Optimizes based on FLOPS (and memory footprint for tie)



# Mathematical property based graph rewriting

- Rewrite candidate search: pattern matching
  - NP-complete problem
- Partition graph on the points where ops don't have mathematical properties
  - Each partition would (*hopefully?*) have tractable search space
  - All cases are considered, and one with least FLOPS is chosen.
- Isn't compiler already doing it after codegen?
  - Operation on the tensors might not be easily optimized as scalars

# Fusion planning: DNN operator classification

Mapping type	Representative
One-to-One	Add, Relu
One-to-Many	Expand
Many-to-Many & N-to-1 (e.g. reduce)	Conv, GEMM
Reorganize	Reshape
Shuffle	Transpose

Second op \ First op	One-to-One	One-to-Many	Many-to-Many	Reorganize	Shuffle
One-to-One	One-to-One	One-to-Many	Many-to-Many	Reorganize	Shuffle
One-to-Many	One-to-Many	One-to-Many	×	One-to-Many	One-to-Many
Many-to-Many	Many-to-Many	Many-to-Many	×	Many-to-Many	Many-to-Many
Reorganize	Reorganize	One-to-Many	Many-to-Many	Reorganize	Reorganize ?
Shuffle	Shuffle	One-to-Many	Many-to-Many	Reorganize	Shuffle

ECG (Extended computational graph) IR will have mapping class annotation for each operator

## Fusion opportunity

**Green:** fusion is profitable

**Yellow(orange?):** further profiling is required (from empirical results)

**Red:** Unprofitable

# Fusion planning: DNN operator classification

Second op \ First op	One-to-One	One-to-Many	Many-to-Many	Reorganize	Shuffle
One-to-One	One-to-One	One-to-Many	Many-to-Many	Reorganize	Shuffle

One-to-one: follows the class of the other fused op

- As the values are directly mapped, fused kernel is not likely to require extra overhead, and only need limited number of registers.

ex) **GEMM(Add(x, y), w)** will be fused to **GEMM(x + y, w)** in codegen step.



# Fusion planning: DNN operator classification

Second op \ First op	One-to-One	One-to-Many	Many-to-Many	Reorganize	Shuffle
Reorganize	Reorganize	One-to-Many	Many-to-Many	Reorganize	Reorganize
Shuffle	Shuffle	One-to-Many	Many-to-Many	Reorganize	Shuffle

Reorder / Shuffle: can be considered as one-to-one with special mapping  
one-to-many / many-to-many fusion should be handled with care

- possible data copying, data access order, redundant computations

ex) **Transpose(Expand(x))** would break continuous memory access pattern when fused

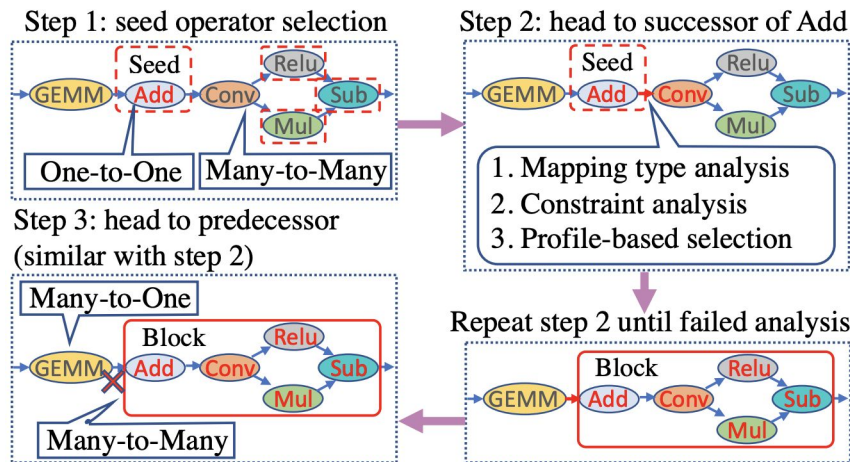
# Fusion planning: DNN operator classification

Second op First op	One-to-Many	Many-to-Many
	One-to-Many	Many-to-Many
One-to-Many	One-to-Many	×
Many-to-Many	Many-to-Many	×

- 1-to-N & N-to-M ( $\text{conv}(\text{Expand}(x))$ ): conv requires **continuous** memory access vs expand would need **distributed** access -> **unprofitable**
- N-to-M & M-to-K ( $\text{conv}(\text{conv}(x))$ ): too complicated. **unprofitable**
- N-to-M & 1-to-K ( $\text{Expand}(\text{conv}(x))$ ): won't affect conv's memory access pattern when expanding one dimension -> **might be OK**
  - $\text{Reshape}(\text{conv}(x))$ : interfere with memory access pattern -> **unprofitable**

# Fusion plan generation

1. Choose **1-to-1** operator with **minimal result** as seed
2. By traversing predecessors and successors, **greedily** group the operators that can be fused
3. Repeat until there's no more fusion candidates



# Other Fusion-related Optimizations

- **Intra**-block optimizations
  - Replace shuffle/reorganize operations with index transform (changed data index)
- **Inter**-block optimizations
  - (Heuristic) Choose one memory layout that would benefit the most **compute-heavy** op in the fusion block

# Codegen

- Generates code for each fused block with data-flow tree (DFT)
- 23 codegen rules, for each mapping class combination
- redundant computation checks from DFT

- Paper says other optimizations were introduced in *PatDNN* (former paper), but not much information
- Seems to have combined all techniques from other frameworks

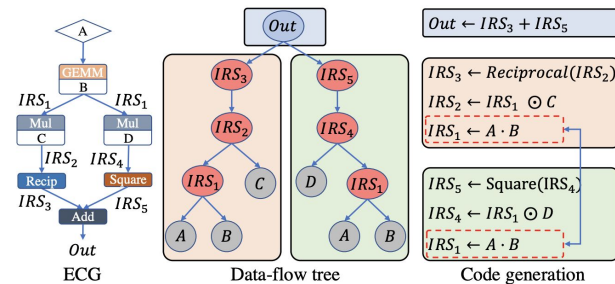


Figure 4. Code generation.

Table 1. DNN acceleration frameworks on mobile devices.

DNNs	Optimization Knobs	TFLite	TVM	MNN	Ours
Dense	Parameters auto-tuning	N	Y	N	Y
	CPU/GPU support	Y	Y	Y	Y
	Half-floating support	Y	Y	Y	Y
	Computation graph optimization	Y <sup>!</sup>	Y <sup>*</sup>	Y <sup>!</sup>	Y <sup>**</sup>
	Tensor optimization	Y <sup>!</sup>	Y <sup>†</sup>	Y <sup>!</sup>	Y <sup>††</sup>

(PatDNN, ASPLOS 2020)

# Putting it all together

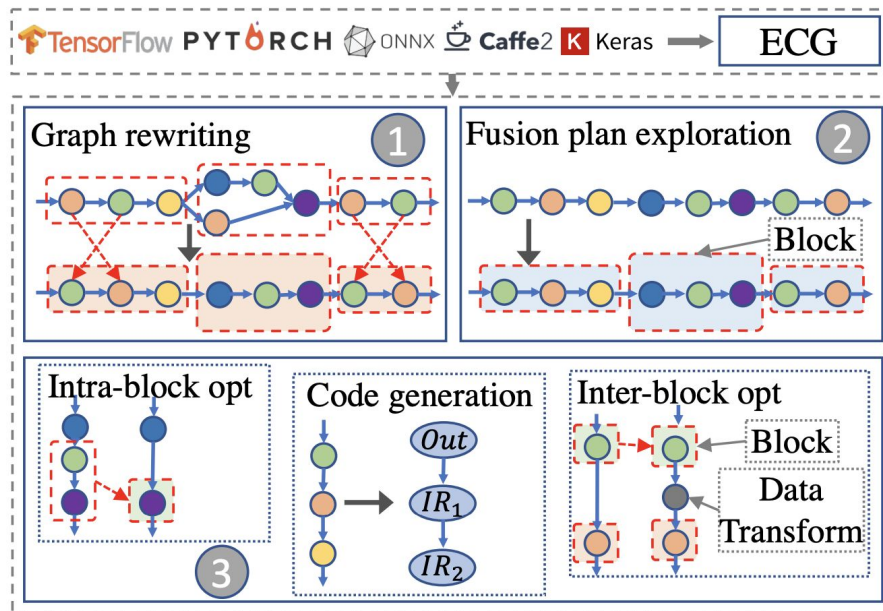


Figure 1. DNNFusion overview.

# Results - fusion

**Table 5. Fusion rate evaluation: computation layer count and intermediate result size for all evaluated DNNs.** CIL (Compute-Intensive Layer): each input is used more than once, e.g. MatMul, CONV. MIL (Memory-Intensive Layer): each input is used only once, e.g. Activation. IRS: intermediate results. '-' means this framework does not support this model.

Model	Type	Task	Layer counts and IRS sizes before opt.				Layer counts and IRS sizes after opt.					
			#CIL	#MIL	#Total layer	IRS size	MNN	TVM	TFLite	Pytorch	DNNF	IRS size
EfficientNet-B0	2D CNN	Image classification	82	227	309	108MB	199	195	201	210	<b>97</b>	<b>26MB</b>
VGG-16	2D CNN	Image classification	16	35	51	161MB	22	22	22	22	<b>17</b>	<b>52MB</b>
MobileNetV1-SSD	2D CNN	Object detection	16	48	202	110MB	138	124	138	148	<b>71</b>	<b>37MB</b>
YOLO-V4	2D CNN	Object detection	106	292	398	329MB	198	192	198	232	<b>135</b>	<b>205MB</b>
C3D	3D CNN	Action recognition	11	16	27	195MB	27	27	-	27	<b>16</b>	<b>90MB</b>
S3D	3D CNN	Action recognition	77	195	272	996MB	-	-	-	272	<b>98</b>	<b>356MB</b>
U-Net	2D CNN	Image segmentation	44	248	292	312MB	241	232	234	-	<b>82</b>	<b>158MB</b>
Faster R-CNN	R-CNN	Image segmentation	177	3,463	3,640	914MB	-	-	-	-	<b>942</b>	<b>374MB</b>
Mask R-CNN	R-CNN	Image segmentation	187	3,812	3,999	1,524MB	-	-	-	-	<b>981</b>	<b>543MB</b>
TinyBERT	Transformer	NLP	37	329	366	183MB	-	304 <sup>†</sup>	322	-	<b>74</b>	<b>55MB</b>
DistilBERT	Transformer	NLP	55	402	457	540MB	-	416 <sup>†</sup>	431	-	<b>109</b>	<b>197MB</b>
ALBERT	Transformer	NLP	98	838	936	1,260MB	-	746 <sup>†</sup>	855	-	<b>225</b>	<b>320MB</b>
BERT <sub>BASE</sub>	Transformer	NLP	109	867	976	915MB	-	760 <sup>†</sup>	873	-	<b>216</b>	<b>196MB</b>
MobileBERT	Transformer	NLP	434	1,953	2,387	744MB	-	1,678 <sup>†</sup>	2,128	-	<b>510</b>	<b>255MB</b>
GPT-2	Transformer	NLP	84	2,449	2,533	1,389MB	-	2,047 <sup>†</sup>	2,223	-	<b>254</b>	<b>356MB</b>

<sup>†</sup> TVM does not support this model on mobile. This layer count number is collected on a laptop platform for reference.

# Results - fusion

Model	Type	Task	Layer counts and IRS sizes before opt.				Layer counts and IRS sizes after opt.					
			#CIL	#MIL	#Total layer	IRS size	MNN	TVM	TFLite	Pytorch	DNNF	IRS size
EfficientNet-B0	2D CNN	Image classification	82	227	309	108MB	199	195	201	210	<b>97</b>	<b>26MB</b>
VGG-16	2D CNN	Image classification	16	35	51	161MB	22	22	22	22	<b>17</b>	<b>52MB</b>
MobileNetV1-SSD	2D CNN	Object detection	16	48	202	110MB	138	124	138	148	<b>71</b>	<b>37MB</b>

Overall impressive, few notes:

- VGG has 22 layers (w/o activation): 13 conv, 5 pooling, 3 dense, 1 softmax
  - 17 seems *13 conv + 3 dense + 1 softmax*. poolings are seems to be fused, although considered many-to-many?
- MobileNetV1-SSD in TFLite can have 64 layers, with predefined fusion of NMS to custom op `TFLite_Detection_PostProcess`.
  - Also shows faster performance than the evaluation result (87ms vs 29ms)
  - Note: 29ms was on Pixel 4, both S20 and Pixel 4 runs on SD865 SoC
  - Hand-crafted fusion kernels still works better for some cases.



# Results - latency improvement

**Table 6. Inference latency comparison: DNNFusion, MNN, TVM, TFLite, and PyTorch on mobile CPU and GPU.** #FLOPS denotes the number of floating point operations. OurB is our baseline implementation by turning off all fusion optimizations and OurB+ is OurB with a fixed-pattern fusion as TVM. DNNF is short for DNNFusion, i.e., our optimized version. '-' denotes this framework does not support this execution.

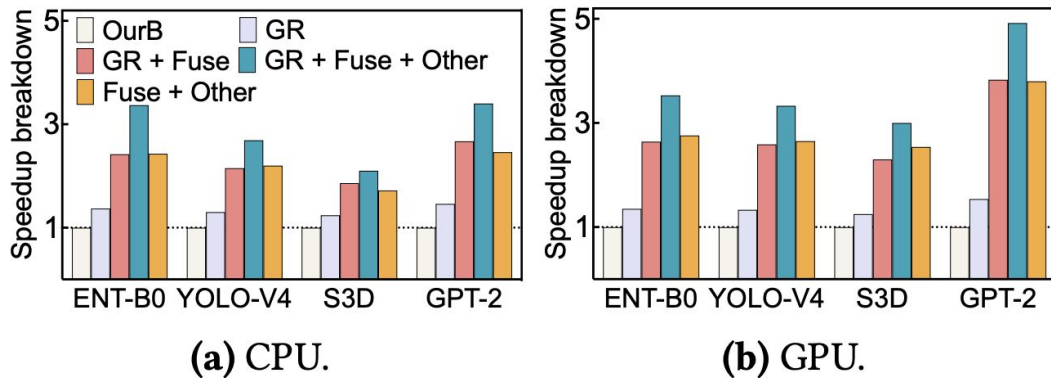
Model	#Params	#FLOPS	MNN (ms)		TVM (ms)		TFLite (ms)		Pytorch (ms)		OurB (ms)		OurB+ (ms)		DNNF (ms)	
			CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU
EfficientNet-B0	5.3M	0.8B	41	26	56	27	52	30	76	-	54	35	38	24	<b>16</b>	<b>10</b>
VGG-16	138M	31.0B	242	109	260	127	245	102	273	-	251	121	231	97	<b>171</b>	<b>65</b>
MobileNetV1-SSD	9.5M	3.0B	67	43	74	52	87	68	92	-	79	56	61	39	<b>33</b>	<b>17</b>
YOLO-V4	64M	34.6B	501	290	549	350	560	288	872	-	633	390	426	257	<b>235</b>	<b>117</b>
C3D	78M	77.0B	867	-	1,487	-	-	-	2,541	-	880	551	802	488	<b>582</b>	<b>301</b>
S3D	8.0M	79.6B	-	-	-	-	-	-	6,612	-	1,409	972	1,279	705	<b>710</b>	<b>324</b>
U-Net	2.1M	15.0B	181	106	210	120	302	117	271	-	227	142	168	92	<b>99</b>	<b>52</b>
Faster R-CNN	41M	47.0B	-	-	-	-	-	-	-	-	2,325	3,054	1,462	1,974	<b>862</b>	<b>531</b>
Mask R-CNN	44M	184B	-	-	-	-	-	-	-	-	5,539	6,483	3,907	4,768	<b>2,471</b>	<b>1,680</b>
TinyBERT	15M	4.1B	-	-	-	-	97	-	-	-	114	89	92	65	<b>51</b>	<b>30</b>
DistilBERT	66M	35.5B	-	-	-	-	510	-	-	-	573	504	467	457	<b>224</b>	<b>148</b>
ALBERT	83M	65.7B	-	-	-	-	974	-	-	-	1,033	1,178	923	973	<b>386</b>	<b>312</b>
BERT <sub>Base</sub>	108M	67.3B	-	-	-	-	985	-	-	-	1,086	1,204	948	1,012	<b>394</b>	<b>293</b>
MobileBERT	25M	17.6B	-	-	-	-	342	-	-	-	448	563	326	397	<b>170</b>	<b>102</b>
GPT-2	125M	69.1B	-	-	-	-	1,102	-	-	-	1,350	1,467	990	1,106	<b>394</b>	<b>292</b>

# Results - latency improvement

**Table 6. Inference latency comparison: DNNFusion, MNN, TVM, TFLite, and PyTorch on mobile CPU and GPU.** #FLOPS denotes the number of floating point operations. OurB is our baseline implementation by turning off all fusion optimizations and OurB+ is OurB with a fixed-pattern fusion as TVM. DNNF is short for DNNFusion, i.e., our optimized version. '-' denotes this framework does not support this execution.

Model	#Params	#FLOPS	MNN (ms)		TVM (ms)		TFLite (ms)		Pytorch (ms)		OurB (ms)		OurB+ (ms)		DNNF (ms)	
			CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU	CPU	GPU
EfficientNet-B0	5.3M	0.8B	41	26	56	27	52	30	76	-	54	35	38	24	<b>16</b>	<b>10</b>
VGG-16	138M	31.0B	242	109	260	127	245	102	273	-	251	121	231	97	<b>171</b>	<b>65</b>
MobileNetV1-SSD	9.5M	3.0B	67	43	74	52	87	68	92	-	79	56	61	39	<b>33</b>	<b>17</b>
YOLO-V4	64M	34.6B	501	290	549	350	560	288	872	-	633	390	426	257	<b>235</b>	<b>117</b>
C3D	78M	77.0B	867	-	1,487	-	-	-	2,541	-	880	551	802	488	<b>582</b>	<b>301</b>
S3D	8.0M	79.6B	-	-	-	-	-	-	6,612	-	1,409	972	1,279	705	<b>710</b>	<b>324</b>
U-Net	2.1M	15.0B	181	106	210	120	302	117	271	-	227	142	168	92	<b>99</b>	<b>52</b>
Faster R-CNN	41M	47.0B	-	-	-	-	-	-	-	-	2,325	3,054	1,462	1,974	<b>862</b>	<b>531</b>
Mask R-CNN	44M	184B	-	-	-	-	-	-	-	-	5,539	6,483	3,907	4,768	<b>2,471</b>	<b>1,680</b>
TinyBERT	15M	4.1B	-	-	-	-	97	-	-	-	114	89	92	65	<b>51</b>	<b>30</b>
DistilBERT	66M	35.5B	-	-	-	-	510	-	-	-	573	504	467	457	<b>224</b>	<b>148</b>
ALBERT	83M	65.7B	-	-	-	-	974	-	-	-	1,033	1,178	923	973	<b>386</b>	<b>312</b>
BERT <sub>Base</sub>	108M	67.3B	-	-	-	-	985	-	-	-	1,086	1,204	948	1,012	<b>394</b>	<b>293</b>
MobileBERT	25M	17.6B	-	-	-	-	342	-	-	-	448	563	326	397	<b>170</b>	<b>102</b>
GPT-2	125M	69.1B	-	-	-	-	1,102	-	-	-	1,350	1,467	990	1,106	<b>394</b>	<b>292</b>


# Results - latency improvement




**Figure 7. Optimization breakdown on y-axis: speedup over OurB, i.e. a version w/o fusion opt.** GR, Fuse, and Other denote graph rewriting, fusion, and other fusion-related optimizations, respectively.

# Results - others

- vs TASO w/ TFLite: 1.4x~2.6x speedup
  - Possibly due to fusion and codegen
- Memory access and cache misses: improved from all frameworks
- G/GPU Utilization: better, > 90%
- Compilation time (vs TVM): best when profiling result can be looked up




**(a)** CPU and GPU utilization.




**(b)** Compilation time.

# Results - others

- vs TASO w/ TFLite: 1.4x~2.6x speedup
  - Possibly due to fusion and codegen
- Memory access and cache misses: improved from all frameworks
- G/GPU Utilization: better, > 90%
- Compilation time (vs TVM): best when profiling result can be looked up
- Can be applied to other targets



(a) CPU and GPU utilization.



(b) Compilation time.

# Takeaways

- Results and related work sections also serve as great **reference points**
- **Simple greedy algorithm** for finding fusion candidates works surprisingly well
- **Unlocked** mobile inference for various models that were not possible for prior frameworks
  
- Accuracy is not reported - might be needed to mention correctness/exactness analysis result after aggressive fusion
- Some details are questionable - fused pooling in VGG, details for codegen