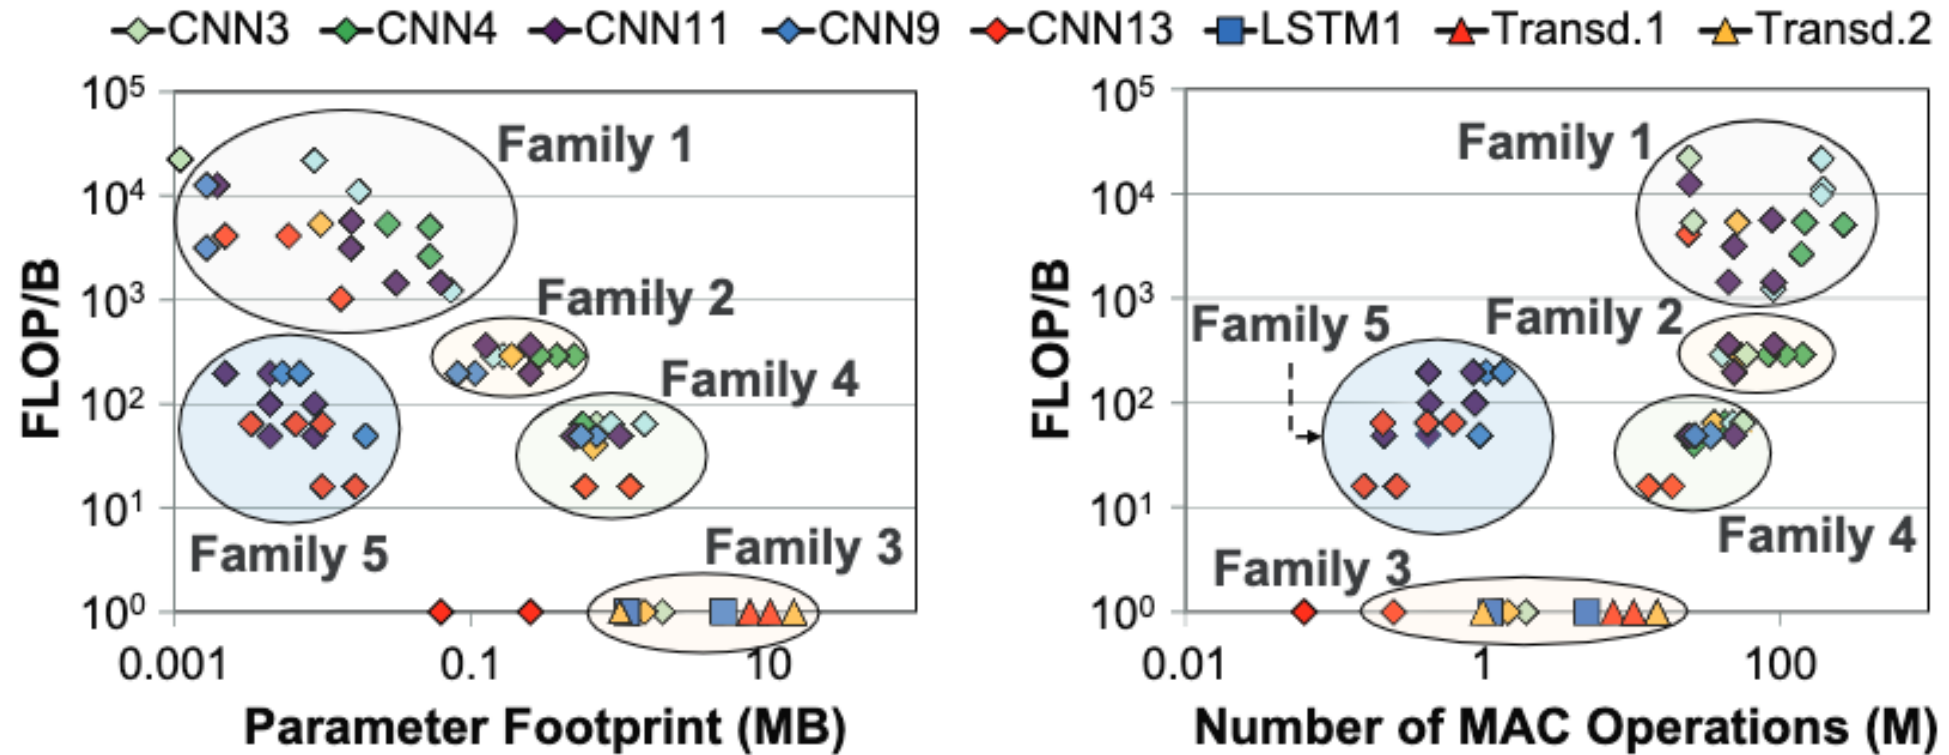


# Google Neural Network Models for Edge Devices: Analyzing and Mitigating Machine Learning Inference Bottleneck

Amirali Boroumand, Saugata Ghose, Berkin Akin Ravi Narayanaswami,  
Geraldo F. Oliveria, Xiaoyu Ma, Eric Shiu, Onur Mutlu

# Model Families



- 5 groups
- Accelerator: F1,2 / F3 / F4,5

# Model Families

- **Family 1**

- Small param. footprint, high AI for params., high MAC AI
- e.g., early layers in CNN/RCNN

- **Family 2**

- small param. footprint, moderate AI for params, high MAC AI
- e.g., pointwise conv, middle

- **Family 3**

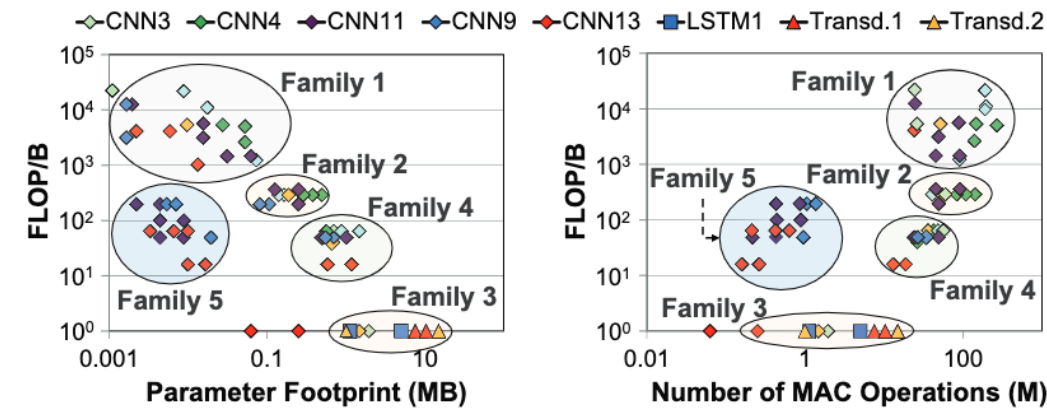
- large param. footprint, minimal AI for params, low MAC AI
- e.g., LSTM, Transducers

- **Family 4**

- relatively large param. footprint, low-to-moderate AI for params, moderate MAC AI
- e.g., deep input/output channels

- **Family 5**

- very small parameter footprint, moderate AI for params, low MAC AI
- e.g., depthwise conv



# Pascal: Compute-Centric Accelerator Design

- Design point
  - Temporal reduction: mitigates the impact of large output activation footprint
  - Should avoid spatial reduction

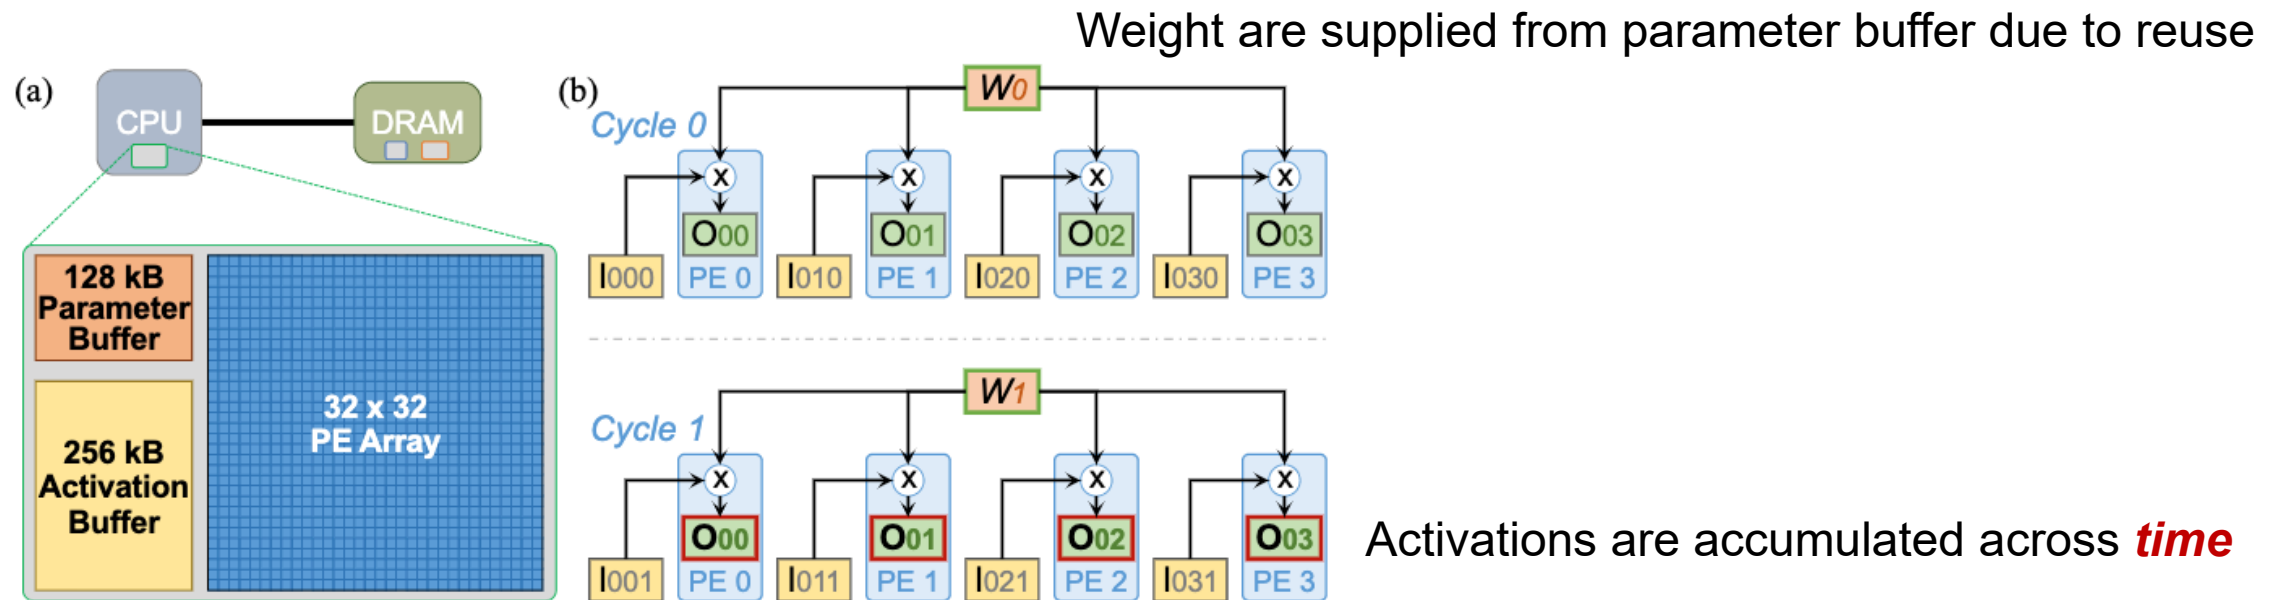


Figure 7. (a) Pascal design; (b) Pascal dataflow for a pointwise layer.

\* Note that Edge TPU has 64x64 PEs and 256kB params. buf.

# Pavlov: LSTM-Centric Accelerator Design

- Design point
  - Exploit output activation reuse opportunities
  - Should reduce the off-chip memory BW required by params.

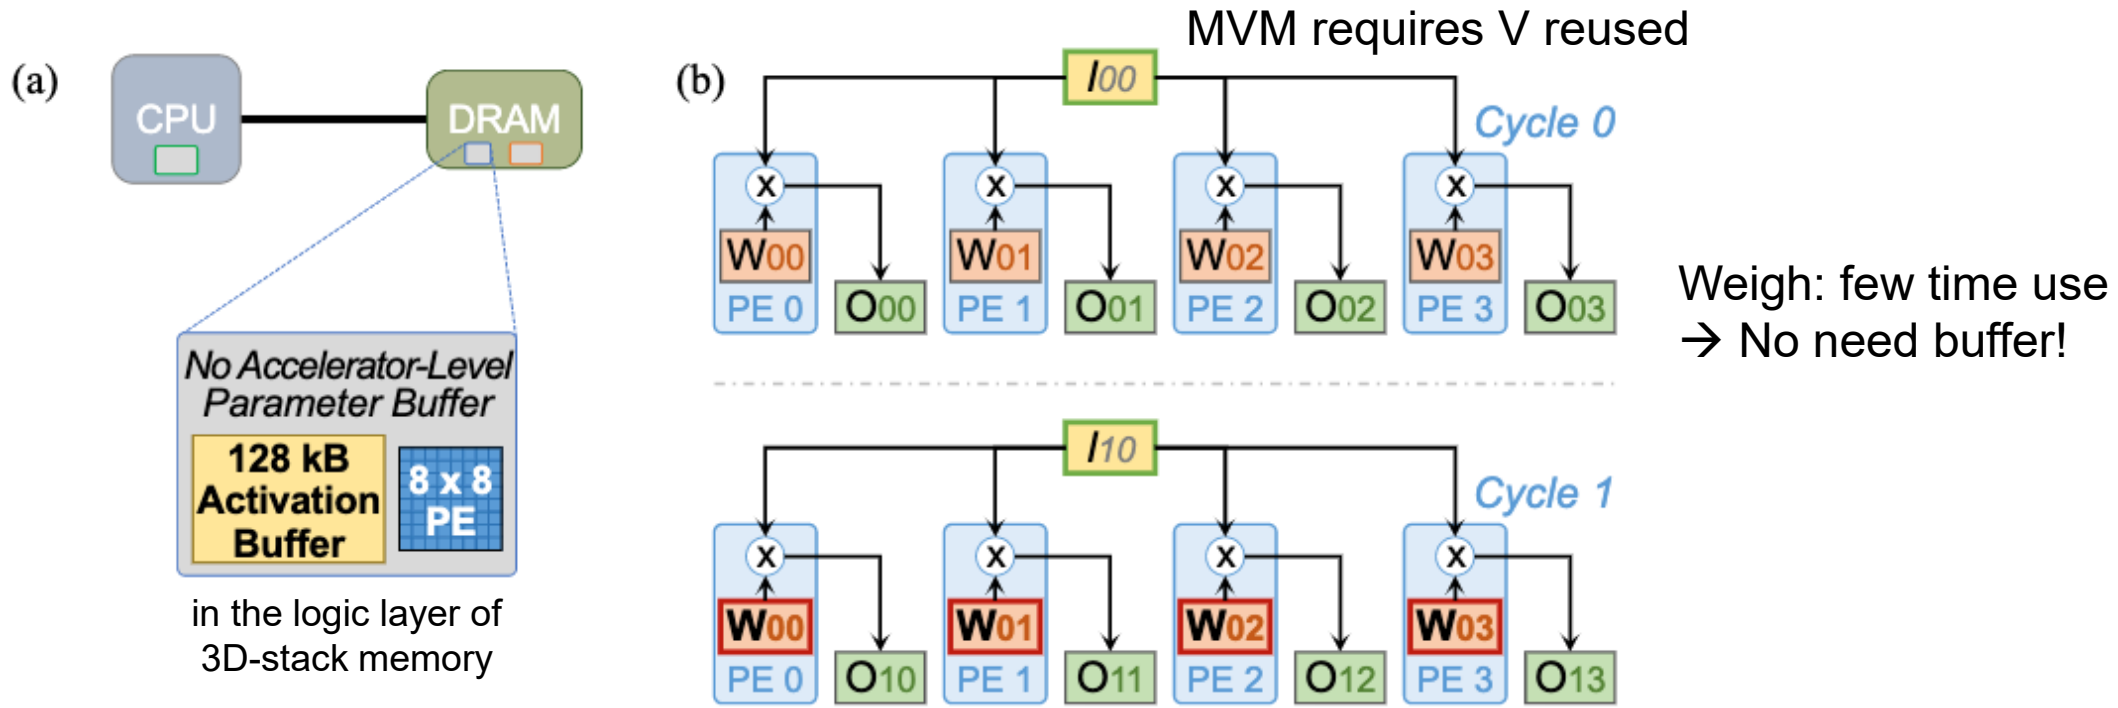
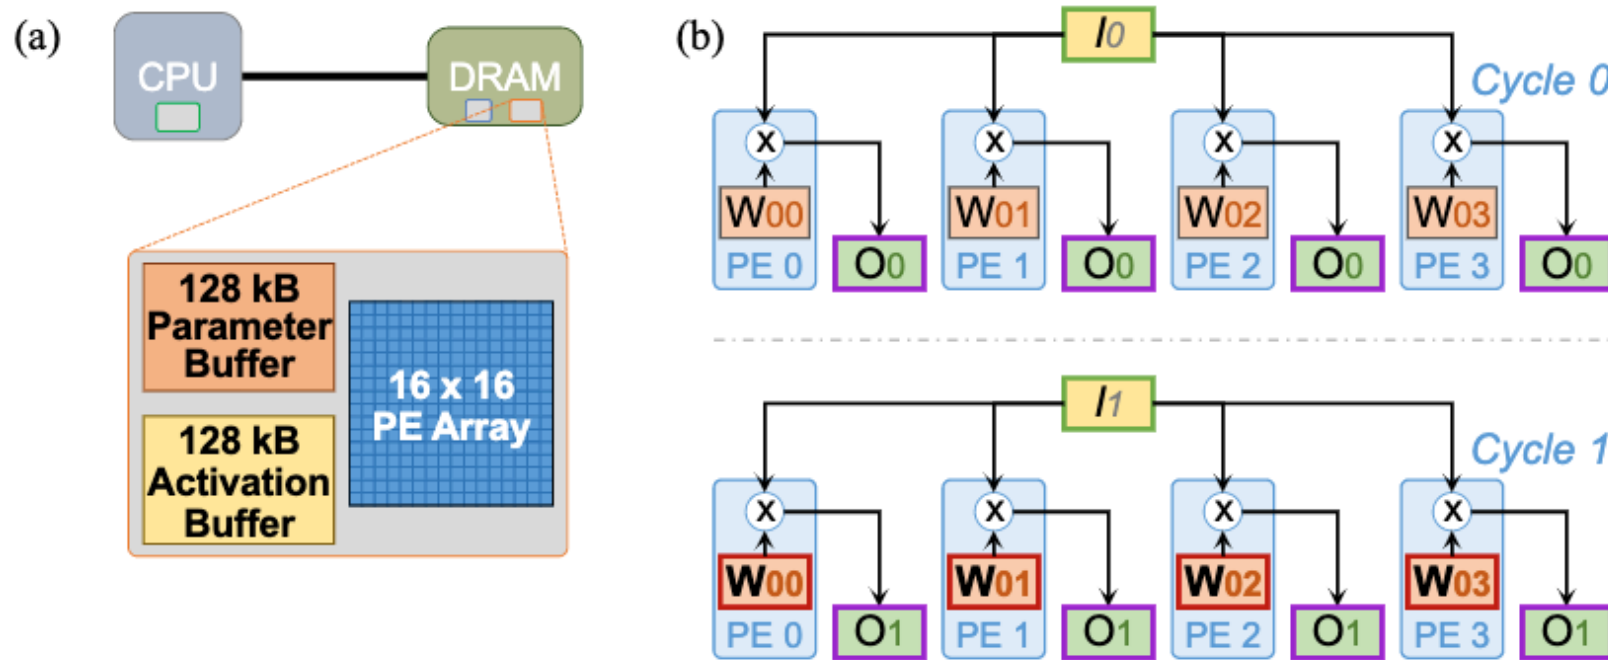


Figure 8. (a) Pavlov design; (b) Pavlov dataflow.

# Jacquard Accelerator Design

- Design point
  - Should exploit temporal reuse opportunities for params.
  - Should provide high off-chip memory bandwidth



Smaller PEs and  
Smaller Act. buf.  
than **Pascal**

But dataflow is similar to Pavlov  
→ due to layer characteristics

Figure 9. (a) Jacquard design; (b) Jacquard dataflow.

# Mensa Runtime Scheduler

- Assumption: No concurrent execution of layers
  - Phase 1
    - Iterates through each layer in the model, and identifies the ideal hardware accelerator *in isolation*
    - To maximize throughput and energy efficiency
    - Not consider transferring act / comm. dependencies.
  - Phase 2
    - Accounts for communication overhead using a simple cost analysis algorithm
    - Two cases
      - If the number of MAC operations required for layer  $i$  is 2x higher than the compute resources available in destination  $i - 1$
      - If the amount of params. data (that destination  $i - 1$  would need to fetch to run layer  $i$ ) is greater than the amount of output activation data (that would have to be sent to the ideal accelerator) AND the opportunities for reusing the params. data in destination  $i - 1$  are low ( $\text{FLOP/B} < 64$ )
- Google edge models typically communicate between accelerators only 4-5 times during execution (implying one-time assignment is mostly optimal)
- ➔ Heuristic-based approach

# Methodology

- Models
  - 24 Google edge NN models from Google mobile applications/products
  - TensorFlow Lite, 8-bit, QAT
  - Not disclosed
- Energy
  - PE and NoC (static/dynamic): In-house simulator
  - On-chip buffer: CACTI-P 6.5
  - MAC: 0.2pJ/bit
  - Memory access: According to LPDDR4 J/bit access in in-house simulator
- Performance
  - In-house simulator (Edge TPU)
  - On-chip buffer latency: CACTI-P 6.5
  - NDP: 256GB/s internal BW of HBM (8x external BW)
  - Edge TPU baseline: uses 2GB of HBM DRAM



# Results




Figure 10. Inference energy across different models (left) and energy breakdown across our three proposed accelerators (right), normalized to Baseline.




Figure 11. PE utilization (top) and Baseline-normalized throughput (bottom) across different models.




Figure 12. Inference latency, normalized to Baseline.

# Interesting Points

- HW/SW co-design
  - Derived from intensive workload analysis
- Heterogeneous system with good scheduler
  - NDP is considered
  - Accelerators with multiple accelerators
- (In-house simulators)
  - Hard to develop new idea on top of this paper
  - Open-sourcing each simulator itself would have good impact

# Future Work

- Concurrent execution and scheduling
  - NP-hard
  - c.f., MAGMA[1]
- Learning-based scheduler
  - Currently hand-tuned
- Heterogeneity accelerator
  - Existing / future modes
  - Possibly more tight co-design of model families – scheduler?
- What about integration with PIM?
  - Enlarging search space with various constraints