

Tensor Comprehensions

Tensor Comprehensions

- ▶ Published in 2018
- ▶ Researcher wants to develop a novel type of layer or network architecture
- ▶ Must develop a customer operator
 - ▶ High engineering cost
 - ▶ Performance penalty
- ▶ Novel domain-specific flow capable of generating highly-optimized kernels for tensor expressions (for GPU)
 - ▶ For new operators, we want to find an efficient implementation easily, without thinking about the hardware

Contributions

- ▶ Domain-specific language “Tensor Comprehensions (TC)”
 - ▶ Syntax is both concise and expressive
 - ▶ Semantics allows for efficient memory management and mapping to complex parallel platforms
- ▶ Specialize a polyhedral intermediate representation and compilation algorithm to DL, provide a dedicated autotuner
- ▶ Enables rapid prototyping of new operators for researchers
- ▶ With comparable performance than manual tuning

Tensor Comprehensions: A Notation

- ▶ Three ideas from Einstein notation
- ▶ Index variables are defined implicitly by using them. Range is inferred.
- ▶ Indices that appear only on the right of an expression are assumed to be reduction dimensions.
- ▶ Evaluation order of points in the iteration space does not affect the output.

```
def mv(float(M,K) A, float(K) x) → (C) {  
    C(i) = 0  
    C(i) += A(i,k) * x(k)  
}
```



```
tensor C({M}).zero(); // 0-filled single-dim tensor  
parallel for (int i = 0; i < M; i++)  
    reduction for (int k = 0; k < K; k++)  
        C(i) += A(i,k) * x(k);
```

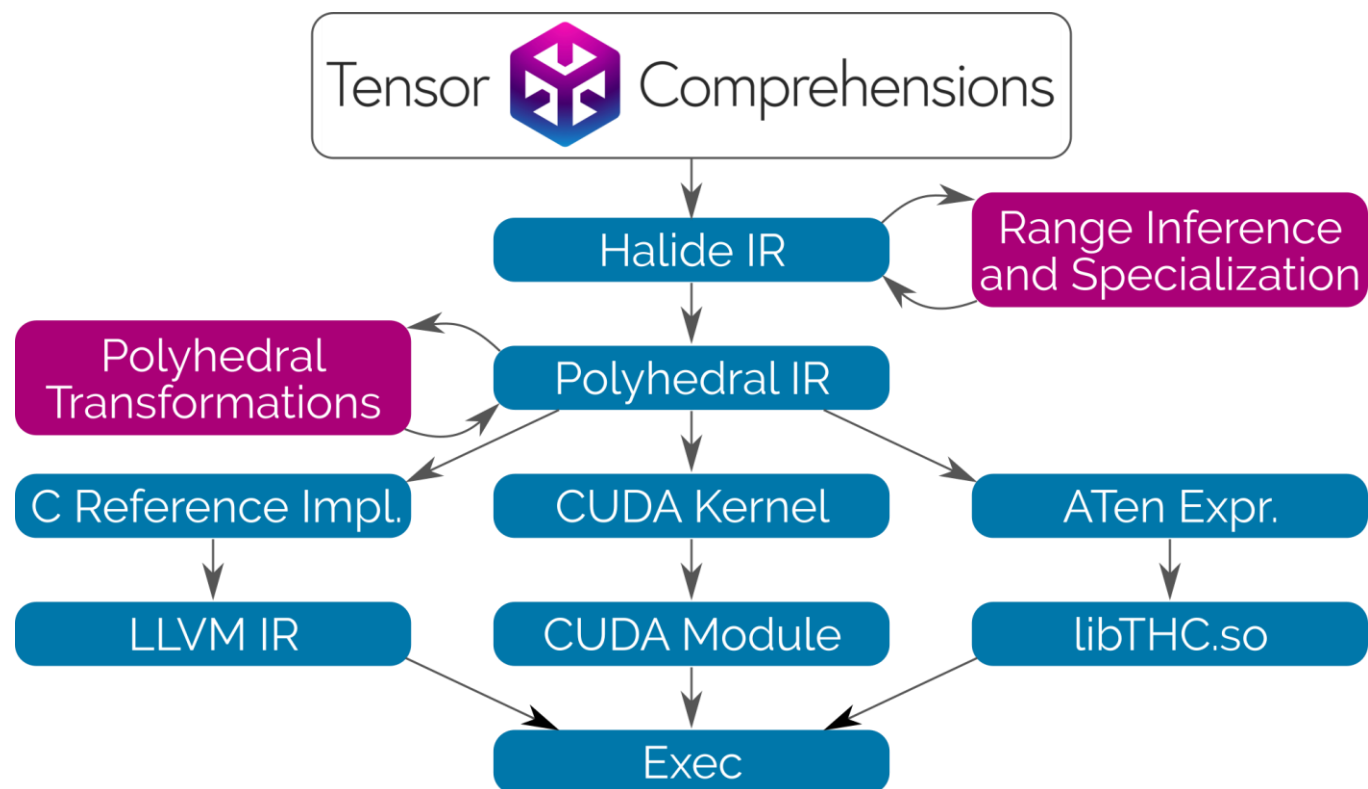
Why TC?

- ▶ TC DSL is extremely simple
- ▶ Resembles the *whiteboard mathematical model* of a deep neural network
- ▶ Makes it easy to reason about, communicate, and to manually alter the computations and storage/computation tradeoffs

```
def sgemm(float a, float b,  
         float(N,M) A, float(M,K) B) → (C) {  
  C(i,j) = b * C(i,j)      # initialization  
  C(i,j) += a * A(i,k) * B(k,j)  # accumulation  
}
```

Figure 1: Tensor Comprehension for the `sgemm` BLAS

TC Flow



TC: Halide + Polyhedral

- ▶ Tensor Comprehensions use Halide and Polyhedral Compilation techniques
 - ▶ Automatically synthesize CUDA kernels
 - ▶ With delegated memory management and synchronization.

Optimization

```
def sgemm(float a, float b,
         float(N,M) A, float(M,K) B) → (C) {
    C(i,j) = b * C(i,j)      # initialization
    C(i,j) += a * A(i,k) * B(k,j)  # accumulation
}
```

Figure 1: Tensor Comprehension for the `sgemm` BLAS

$$C_{i,j} = \sum_{k=1}^N A_{i,k} \cdot B_{k,j}, \quad \forall i, j \in 1, N$$

Domain $\left[\begin{array}{l} \{S(i, j) \mid 0 \leq i < N \wedge 0 \leq j < K\} \\ \{T(i, j, k) \mid 0 \leq i < N \\ \wedge 0 \leq j < K \wedge 0 \leq k < M\} \end{array} \right.$

Sequence

Filter{S(i, j)}

Band{S(i, j) → (i, j)}

Filter{T(i, j, k)}

Band{T(i, j, k) → (i, j, k)}

(a) canonical `sgemm`

Domain $\left[\begin{array}{l} \{S(i, j) \mid 0 \leq i < N \wedge 0 \leq j < K\} \\ \{T(i, j, k) \mid 0 \leq i < N \wedge 0 \leq j < K \wedge 0 \leq k < M\} \end{array} \right.$

Band $\left[\begin{array}{l} \{S(i, j) \rightarrow (i, j)\} \\ \{T(i, j, k) \rightarrow (i, j)\} \end{array} \right.$

Sequence

Filter{S(i, j)}

Filter{T(i, j, k)}

Band{T(i, j, k) → (k)}

(b) fused

Domain $\left[\begin{array}{l} \{S(i, j) \mid 0 \leq i < N \wedge 0 \leq j < K\} \\ \{T(i, j, k) \mid 0 \leq i < N \\ \wedge 0 \leq j < K \wedge 0 \leq k < M\} \end{array} \right.$

Band $\left[\begin{array}{l} \{S(i, j) \rightarrow (32[i/32], 32[j/32])\} \\ \{T(i, j, k) \rightarrow (32[i/32], 32[j/32])\} \end{array} \right.$

Band $\left[\begin{array}{l} \{S(i, j) \rightarrow (i \bmod 32, j \bmod 32)\} \\ \{T(i, j, k) \rightarrow (i \bmod 32, j \bmod 32)\} \end{array} \right.$

Sequence

Filter{S(i, j)}

Filter{T(i, j, k)}

Band{T(i, j, k) → (k)}

(c) fused and tiled

Domain $\left[\begin{array}{l} \{S(i, j) \mid 0 \leq i < N \wedge 0 \leq j < K\} \\ \{T(i, j, k) \mid 0 \leq i < N \wedge 0 \leq j < K \wedge 0 \leq k < M\} \end{array} \right.$

Band $\left[\begin{array}{l} \{S(i, j) \rightarrow (32[i/32], 32[j/32])\} \\ \{T(i, j, k) \rightarrow (32[i/32], 32[j/32])\} \end{array} \right.$

Sequence

Filter{S(i, j)}

Band{S(i, j) → (i mod 32, j mod 32)}

Filter{T(i, j, k)}

Band{T(i, j, k) → (32[k/32])}

Band{T(i, j, k) → (k mod 32)}

Band{T(i, j, k) → (i mod 32, j mod 32)}

(d) fused, tiled and sunk

Domain $\left[\begin{array}{l} \{S(i, j) \mid 0 \leq i < N \wedge 0 \leq j < K\} \\ \{T(i, j, k) \mid 0 \leq i < N \wedge 0 \leq j < K \wedge 0 \leq k < M\} \end{array} \right.$

Context $\{0 \leq b_x, b_y < 32 \wedge 0 \leq t_x, t_y < 16\}$

Filter $\left[\begin{array}{l} \{S(i, j) \mid i - 32b_x - 31 \leq 32 \times 16[i/32/16] \leq i - 32b_x \wedge \\ j - 32b_y - 31 \leq 32 \times 16[j/32/16] \leq j - 32b_y\} \\ \{T(i, j, k) \mid i - 32b_x - 31 \leq 32 \times 16[i/32/16] \leq i - 32b_x \wedge \\ j - 32b_y - 31 \leq 32 \times 16[j/32/16] \leq j - 32b_y\} \end{array} \right.$

Band $\left[\begin{array}{l} \{S(i, j) \rightarrow (32[i/32], 32[j/32])\} \\ \{T(i, j, k) \rightarrow (32[i/32], 32[j/32])\} \end{array} \right.$

Sequence

Filter{S(i, j)}

Filter{S(i, j) | (t_x - i) = 0 mod 16 ∧ (t_y - j) = 0 mod 16}

Band{S(i, j) → (i mod 32, j mod 32)}

Filter{T(i, j, k)}

Band{T(i, j, k) → (32[k/32])}

Band{T(i, j, k) → (k mod 32)}


Filter{T(i, j, k) | (t_x - i) = 0 mod 16 ∧ (t_y - j) = 0 mod 16}

Band{T(i, j, k) → (i mod 32, j mod 32)}


(e) fused, tiled, sunk and mapped

Optimization

Affine Loop Transformations




Loop Fusion and Fission



$$A(i, j) = 42 * X(j)$$
$$B(j, i) = B(j, i) + A(i, j)$$


TC: Polyhedral Transformaion + Mapping

- Transform for parallelism and data locality
- Map GPU compute and memory resources to the transformed program



$$\begin{bmatrix} 1 & 0 \\ -1 & 0 \\ 0 & 1 \\ 1 & -1 \end{bmatrix} \begin{pmatrix} i \\ j \end{pmatrix} \geq \begin{pmatrix} 1 \\ -n \\ 1 \\ 1 \end{pmatrix}$$

(a) original polyhedron $A\vec{x} \geq \vec{c}$



$$\begin{bmatrix} 1/2 & -1/2 \\ -1/2 & 1/2 \\ 0 & 1 \\ 1/2 & -3/2 \end{bmatrix} \begin{pmatrix} i' \\ j' \end{pmatrix} \geq \begin{pmatrix} 1 \\ -n \\ 1 \\ 1 \end{pmatrix}$$

(b) usual transformation $(AT^{-1})\vec{y} \geq \vec{c}$

Domain $\left[\begin{array}{l} \{S(i, j) \mid 0 \leq i < N \wedge 0 \leq j < K\} \\ \{T(i, j, k) \mid 0 \leq i < N \wedge 0 \leq j < K \wedge 0 \leq k < M\} \end{array} \right.$

Context $\{0 \leq b_x, b_y < 32 \wedge 0 \leq t_x, t_y < 16\}$

Filter $\left[\begin{array}{l} \{S(i, j) \mid i - 32b_x - 31 \leq 32 \times 16 \lfloor i/32/16 \rfloor \leq i - 32b_x \wedge \\ j - 32b_y - 31 \leq 32 \times 16 \lfloor j/32/16 \rfloor \leq j - 32b_y\} \\ \{T(i, j, k) \mid i - 32b_x - 31 \leq 32 \times 16 \lfloor i/32/16 \rfloor \leq i - 32b_x \wedge \\ j - 32b_y - 31 \leq 32 \times 16 \lfloor j/32/16 \rfloor \leq j - 32b_y\} \end{array} \right.$

Band $\left[\begin{array}{l} \{S(i, j) \rightarrow (32 \lfloor i/32 \rfloor, 32 \lfloor j/32 \rfloor)\} \\ \{T(i, j, k) \rightarrow (32 \lfloor i/32 \rfloor, 32 \lfloor j/32 \rfloor)\} \end{array} \right.$

Sequence

Filter $\{S(i, j)\}$

Filter $\{S(i, j) \mid (t_x - i) = 0 \bmod 16 \wedge (t_y - j) = 0 \bmod 16\}$

Band $\{S(i, j) \rightarrow (i \bmod 32, j \bmod 32)\}$

Filter $\{T(i, j, k)\}$

Band $\{T(i, j, k) \rightarrow (32 \lfloor k/32 \rfloor)\}$

Band $\{T(i, j, k) \rightarrow (k \bmod 32)\}$

Filter $\{T(i, j, k) \mid (t_x - i) = 0 \bmod 16 \wedge (t_y - j) = 0 \bmod 16\}$

Band $\{T(i, j, k) \rightarrow (i \bmod 32, j \bmod 32)\}$

(e) fused, tiled, sunk and mapped

Autotuning

- ▶ Autotuner interacts with the rest of the environment through the compilation cache: best versions are stored for later use.
- ▶ Compilation cache stores the generated CUDA or PTX code for a given TC
 - ▶ Generated code depends on the input shapes, the selected optimization options, constraints induced by the target GPU architecture
- ▶ Autotuner runs for a prescribed amount of time, updating the cache with better versions along the way using genetic algorithm
 - ▶ Three parents are selected probabilistically based on their fitness, the higher the fitness the higher the selection chance
 - ▶ Each “gene”, which corresponds to one tuning parameter, of the new candidate is randomly selected from the parents.

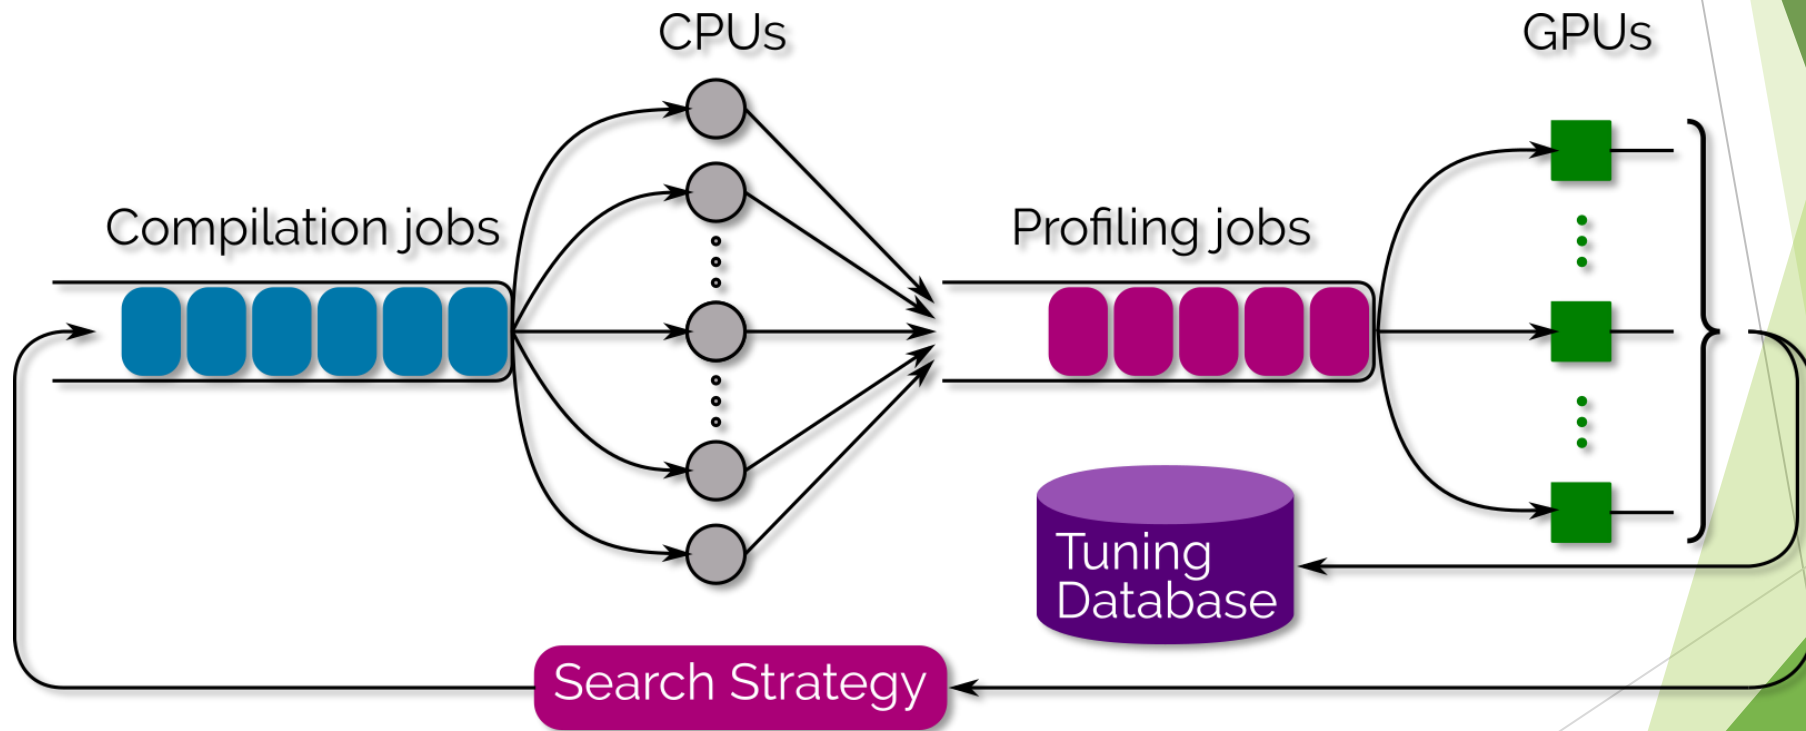
Autotuning

```
def avgpool(float(B, C, H, W) input) -> (output) {{  
    output(b, c, h, w) += input(b, c, h * {sH} + kh, w * {sW} + kw)  
    where kh in 0:{kH}, kw in 0:{kW}  
}}
```




Tensor Comprehension for 2D Average Pooling

Multithreaded autotuning pipeline



TC Performance



```

import tensor_comprehensions as tc
import torch
lang = """
def tensordot(float(N, C1, C2, H, W) I0, float(N, C2, C3, H, W) I1) -> (O) {
    O(n, c1, c3, h, w) +=! I0(n, c1, c2, h, w) * I1(n, c2, c3, h, w)
}
"""

N, C1, C2, C3, H, W = 32, 512, 8, 2, 28, 28
tensordot = tc.define(lang, name="tensordot")
I0, I1 = torch.randn(N, C1, C2, H, W).cuda(), torch.randn(N, C2, C3, H, W).cuda()
best_options = tensordot.autotune(I0, I1, cache=True)
out = tensordot(I0, I1, options=best_options)

```

```

ntv@devfair0172:~/tensorComprehensions$ ./build/test/test_autotuner --smoke_check=0 --tuner_threads=20 --tuner_gpus="0,1" --gtest_filter="*Dot*"
Note: Google Test filter = *Dot*
[=====] Running 1 test from 1 test case.
[-----] Global test environment set-up.
[-----] 1 test from ATenCompilationUnitTest
[ RUN     ] ATenCompilationUnitTest.TensorDot

-----
----- KERNEL STATS -----
----- 100 ITERATIONS -----
-----
Min: 4881us, p50: 4936us, p90: 5134us, p99: 5403us, Max: 5403us
-----

----- TOTAL STATS -----
----- 100 ITERATIONS -----
-----
Min: 4903us, p50: 4947us, p90: 5138us, p99: 5375us, Max: 5375us
-----

Generation 0  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 4177/11678/621574
Generation 1  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2986/6414/20158
Generation 2  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2270/6193/14676
Generation 3  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2267/5608/11261
Generation 4  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2266/4817/11330
Generation 5  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2258/4632/11264
Generation 6  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2242/4634/10955
Generation 7  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2247/4488/10950
Generation 8  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2255/4333/10948
Generation 9  Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2253/4259/10370
Generation 10 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2233/4261/10364
Generation 11 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2233/4171/10371
Generation 12 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 1968/4198/10365
Generation 13 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 1959/3904/8670
Generation 14 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 2084/3426/7846
Generation 15 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 1959/3483/8929
Generation 16 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 1959/3651/9465
Generation 17 Jobs(Compiled, GPU)/total (100, 100)/100 (best/median/worst)us: 1812/3636/9392
Generation 18 Jobs(Compiled, GPU)/total (100, 99)/100 (best/median/worst)us: 1784/3409/8828

```